

# Lab 10: Regression Tress, Random Forest (Problem 37)

AUTHOR

Aaron Younger

## Business Understanding

Jerry Stevenson is the manager of a travel agency. He wants to build a model that can predict customers' annual spending on travel products. A Regression Tree and the Ensemble Method random forest will be used for predicting the customer's annual household spending on travel products.

## Data Understanding

### R Version

```
options(scipen=999)
suppressWarnings(RNGversion("3.5.3"))
```

## Libraries

```
library(readxl)
library(DataExplorer)
library(tidyverse)
library(dplyr)
library(ggplot2)
library(dlookr)
library(e1071)
library(caret)
library(rpart)
library(rpart.plot)
library(randomForest)
library(pdp)
library(gridExtra)
```

```
library(readxl)
travel_data <- read_excel("jaggia_ba_2e_ch13_data.xlsx",
    sheet = "Travel_Data")
View(travel_data)

library(readxl)
travel_score <- read_excel("jaggia_ba_2e_ch13_data.xlsx",
    sheet = "Travel_Score")
View(travel_score)
```

### Data Set Key:

- College: A binary variable representing if the individual went to College.
- Credit Card: A binary variable representing if the individual has a credit card.
- Food Spend: Annual household spending on Food.
- Income: The individuals annual income.
- Travel Spend: The annual household spending on travel products. The dependent variable.

## Dataset Exploration

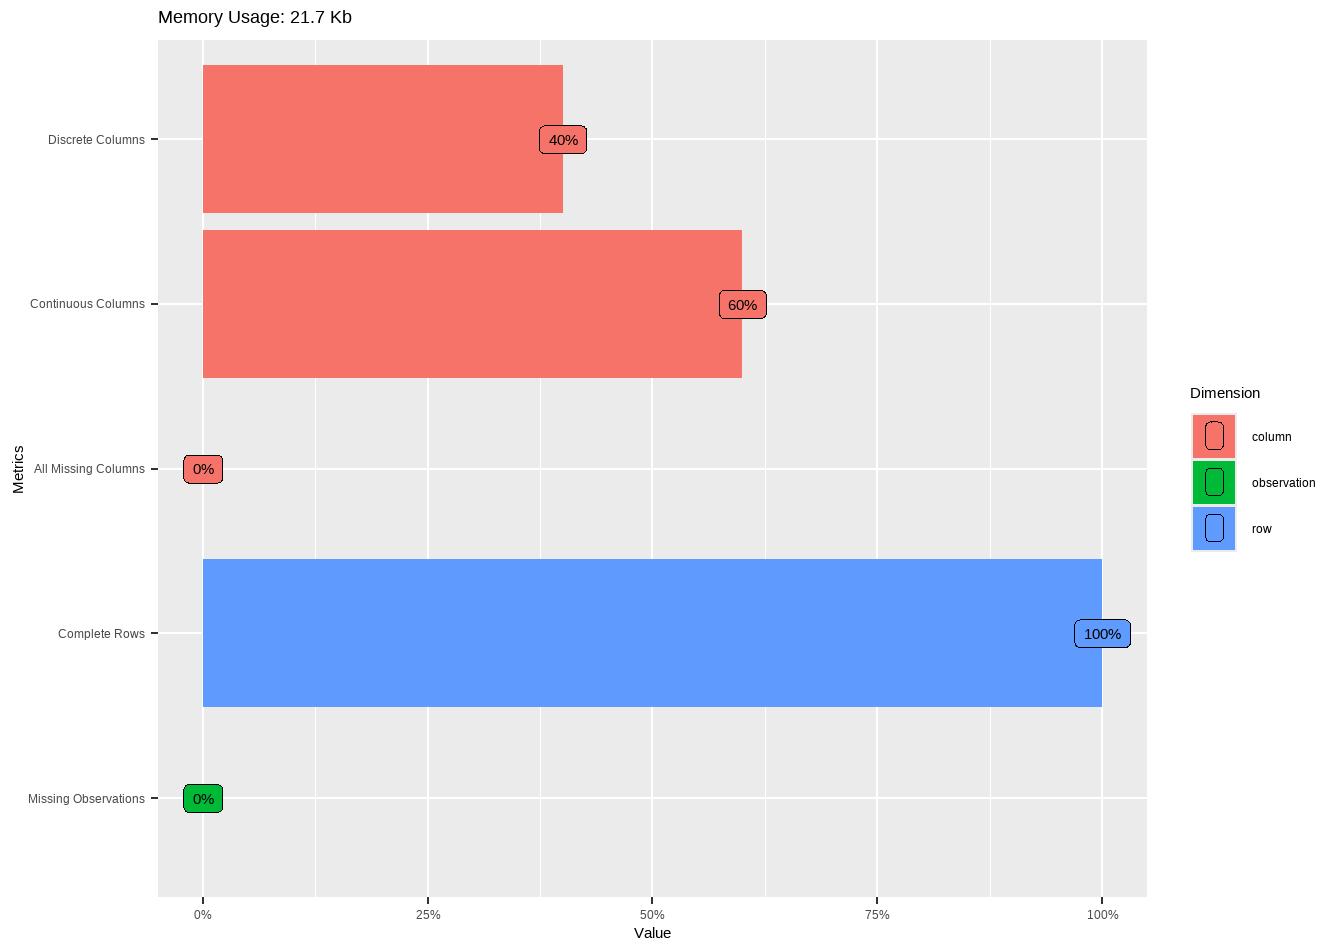
```
travel_data %>% head()
```

```
# A tibble: 6 × 5
  College CreditCard FoodSpend Income TravelSpend
  <chr>   <chr>     <dbl>    <dbl>      <dbl>
1 Yes     Yes        5472.   49150      827.
2 No      Yes        9131.   47806      864.
3 No      Yes        4451.   46050     1180.
4 No      Yes        5021.   42600      756.
5 No      Yes        3408.   52092      851.
6 No      No         2850.   60248     1752.
```

```
travel_data %>% tail()
```

```
# A tibble: 6 × 5
  College CreditCard FoodSpend Income TravelSpend
  <chr>   <chr>     <dbl>    <dbl>      <dbl>
1 No      No         4711.   57274      3448.
2 Yes     No         3787.   65982      2059.
3 No      No         6461.   57274      3853.
4 No      No         3586.   42250      2252.
5 Yes     No         1834.   60248      2000.
6 No      No         6206.   58520      3668.
```

```
travel_data %>% plot_intro()
```

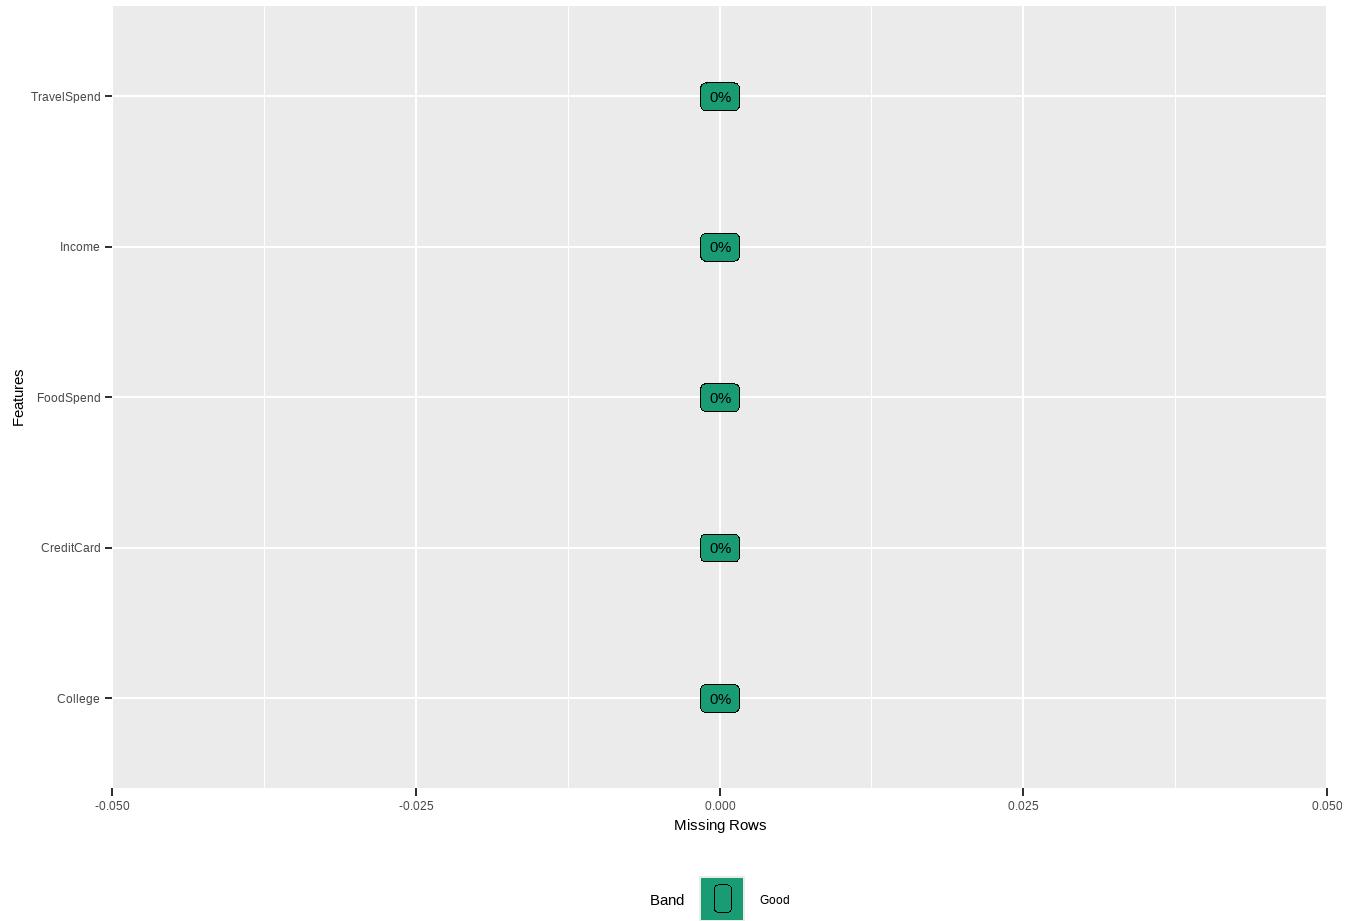


```
travel_data %>% glimpse()
```

Rows: 500  
 Columns: 5

```
$ College      <chr> "Yes", "No", "No", "No", "No", "No", "No", "No", "No"...
$ CreditCard   <chr> "Yes", "Yes", "Yes", "Yes", "No", "No", "Yes", "No"...
$ FoodSpend    <dbl> 5472.43, 9130.73, 4450.67, 5020.72, 3408.11, 2850.33, 3680...
$ Income       <dbl> 49150, 47806, 46050, 42600, 52092, 60248, 49412, 58520, 57...
$ TravelSpend  <dbl> 827.40, 863.55, 1180.05, 755.70, 851.40, 1751.70, 4557.30,...
```

```
travel_data %>% plot_missing()
```

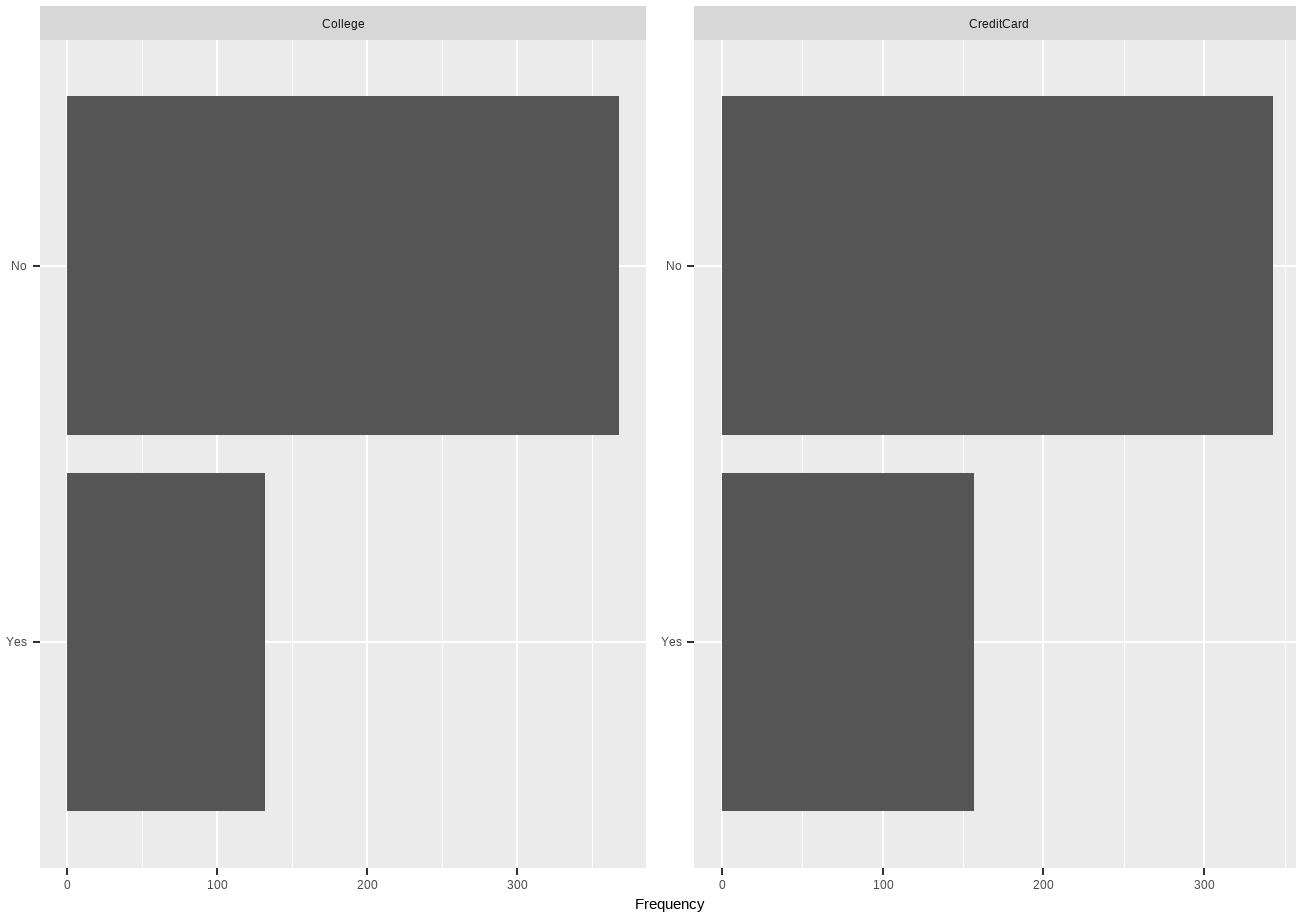


```
travel_data %>% str()
```

```
tibble [500 x 5] (S3: tbl_df/tbl/data.frame)
$ College    : chr [1:500] "Yes" "No" "No" "No" ...
$ CreditCard : chr [1:500] "Yes" "Yes" "Yes" "Yes" ...
$ FoodSpend   : num [1:500] 5472 9131 4451 5021 3408 ...
$ Income      : num [1:500] 49150 47806 46050 42600 52092 ...
$ TravelSpend: num [1:500] 827 864 1180 756 851 ...
```

## Variable EDA

```
travel_data %>% plot_bar()
```



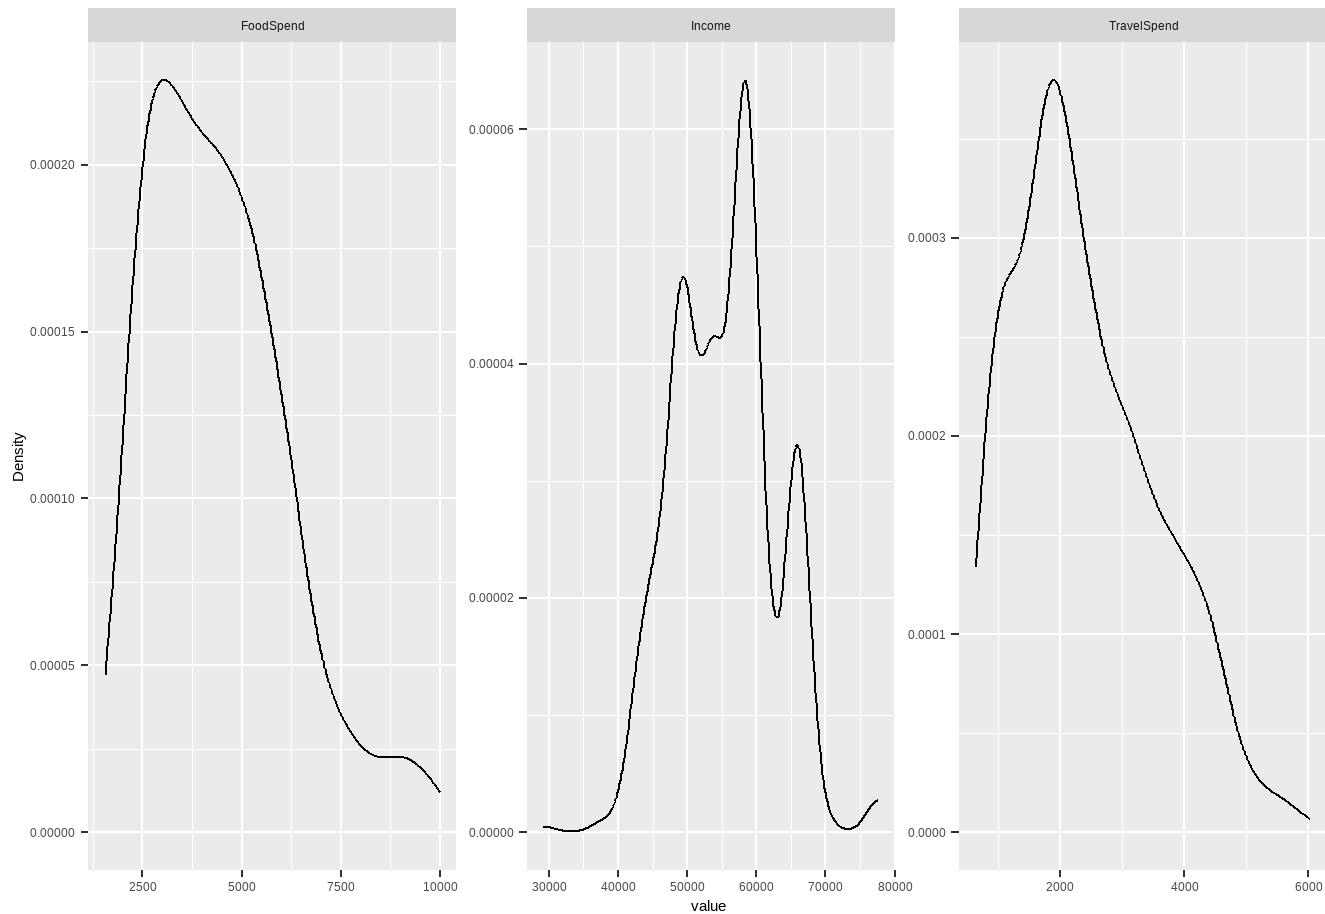
```
prop.table(table(travel_data$College))
```

No	Yes
0.736	0.264

```
prop.table(table(travel_data$CreditCard)) ## Categorical Variables are disproportional, dominant level is the "no" value for both variables.
```

No	Yes
0.686	0.314

```
travel_data %>% plot_density()
```



```
skew_table <- data.frame(
  Variable = c("Food Spend", "Income", "Travel Spend"),
  Skewness = c(
    skewness(travel_data$FoodSpend),
    skewness(travel_data$Income),
    skewness(travel_data$TravelSpend)
  )
)

print(skew_table)
```

	Variable	Skewness
1	Food Spend	0.8343613
2	Income	0.1607550
3	Travel Spend	0.6157839

#### Comments on Variable EDA:

Both Categorical Variables are Binary and have observations in both levels. However, the levels are disproportionate with the dominant level being “no” for both College and Credit Card. The density plots for the numeric variables show that none of the variables are normally distributed. Food and Travel Spend show mild right skewness where most families spend lower on money on food and travel with a few families that spend higher amounts, pulling the distribution to the right. The Income variable shows some multimodality, which suggests possible subgroups within the data. The skewness values for these numeric variables are

positive and less than one, supporting right skewness but suggesting that these variables can be considered approximately normal.

## Outliers

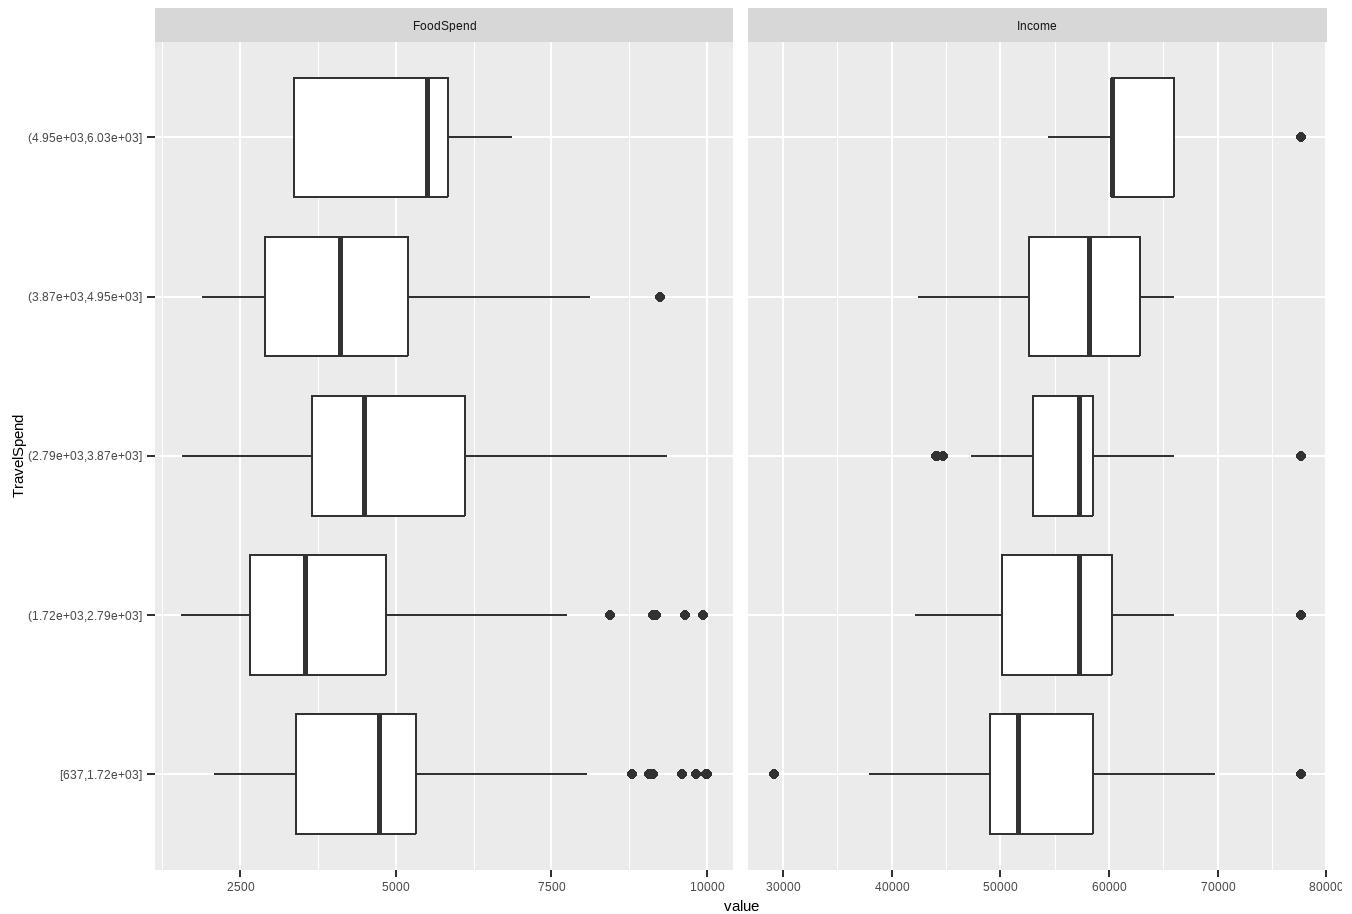
```
diagnose_outlier(travel_data)
```

```
# A tibble: 3 × 6
  variables outliers_cnt outliers_ratio outliers_mean with_mean without_mean
  <chr>        <int>       <dbl>      <dbl>      <dbl>      <dbl>
1 FoodSpend     12          2.4       9516.     4420.     4295.
2 Income         7          1.4      70708      55416.    55199.
3 TravelSpend   5           1       5679.     2402.     2369.
```

```
diagnose_outlier(travel_score)
```

```
# A tibble: 2 × 6
  variables outliers_cnt outliers_ratio outliers_mean with_mean without_mean
  <chr>        <int>       <dbl>      <dbl>      <dbl>      <dbl>
1 FoodSpend     0           0        NaN       4455.     4455.
2 Income         0           0        NaN      59984     59984
```

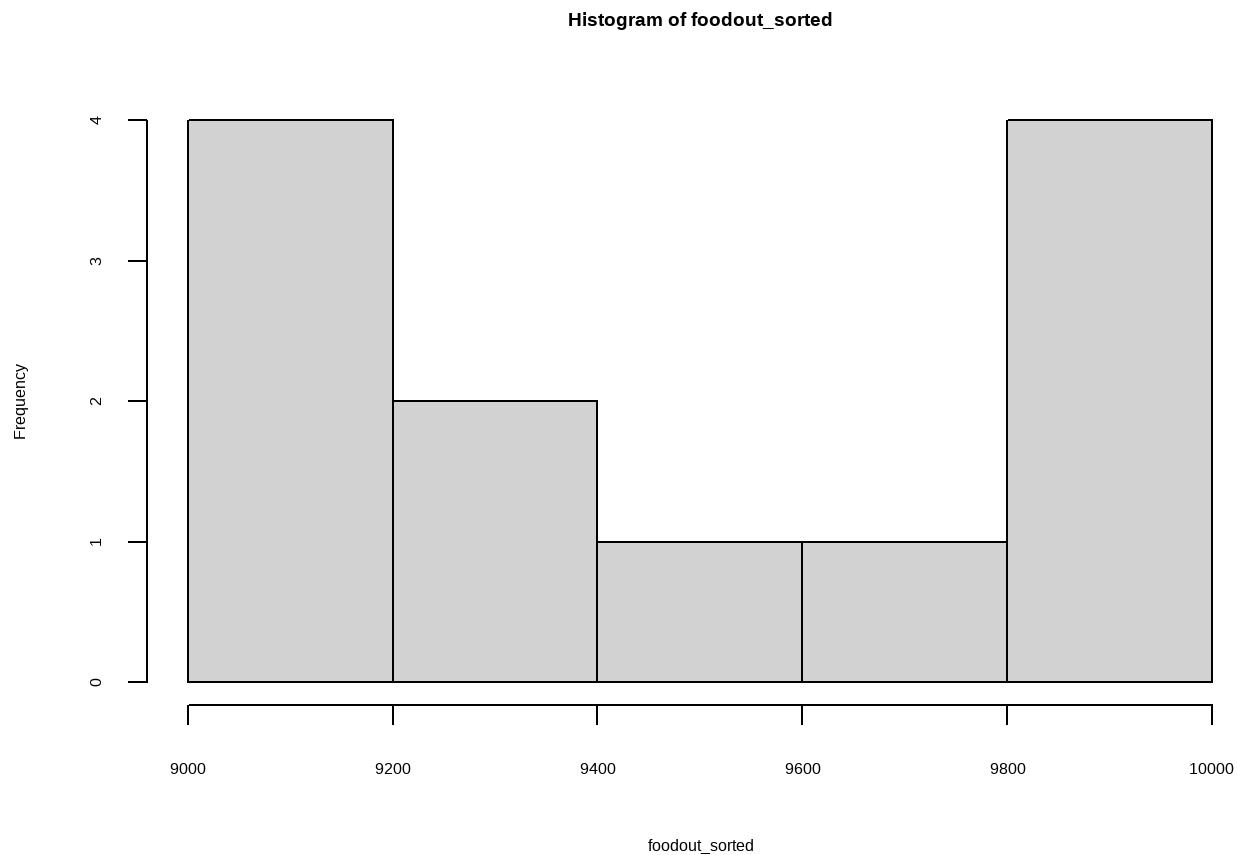
```
travel_data %>% plot_boxplot(by="TravelSpend")
```



```
foodout <- boxplot.stats(travel_data$FoodSpend)$out
foodout_sorted <- sort(foodout, decreasing = TRUE)
foodout_sorted
```

```
[1] 10000.00 9978.49 9935.07 9819.56 9649.01 9592.99 9350.13 9249.13
[9] 9185.28 9174.83 9130.73 9129.66
```

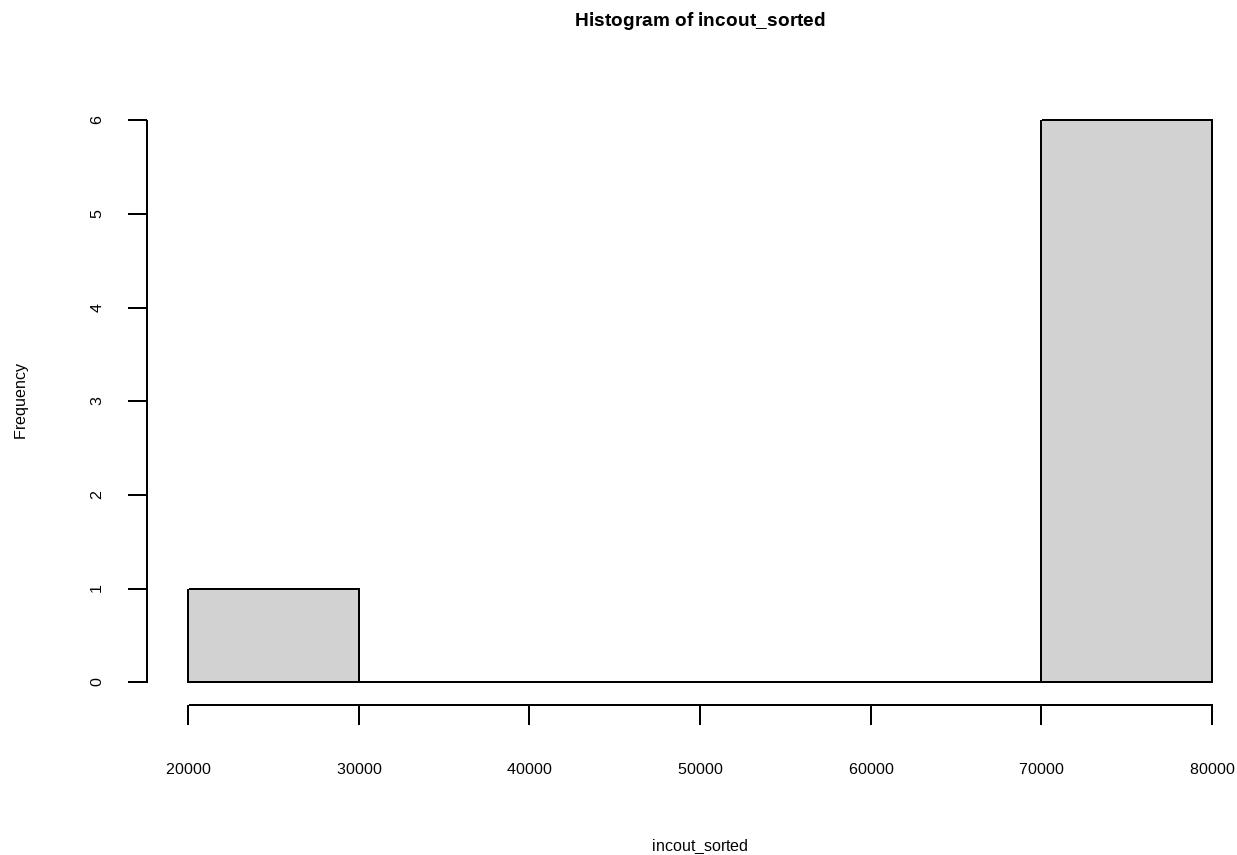
```
hist(foodout_sorted)
```



```
incout <- boxplot.stats(travel_data$Income)$out  
incout_sorted <- sort(incout, decreasing = TRUE)  
incout_sorted
```

```
[1] 77626 77626 77626 77626 77626 77626 77626 29200
```

```
hist(incout_sorted)
```



#### Comments on Outliers:

There are no outliers in the score dataset. There are outliers in the modeling dataset but are not extreme enough to delete.

## Mean of Dependent Variable

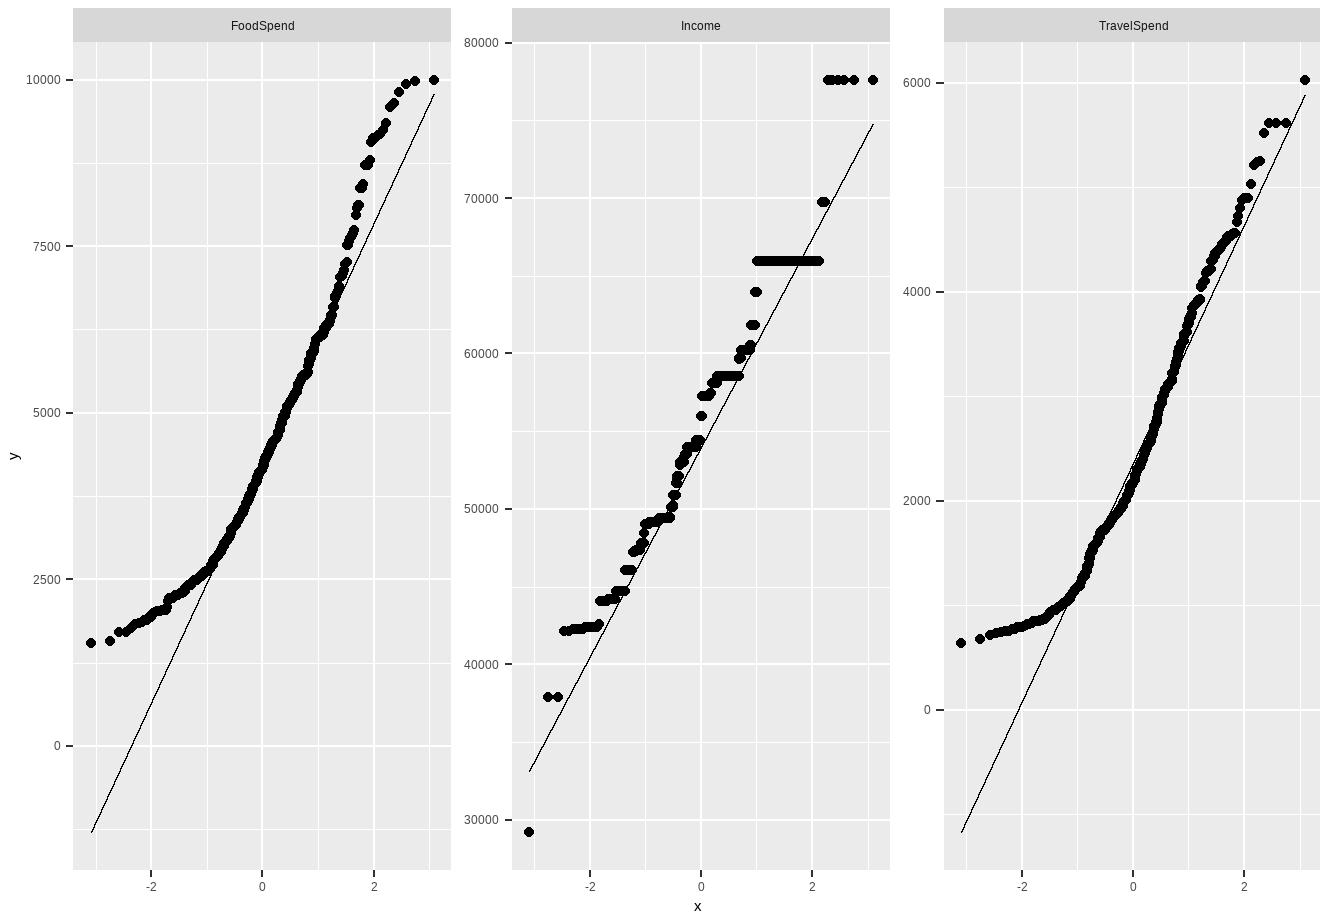
```
mean(travel_data$TravelSpend)
```

```
[1] 2402.287
```

The Average value of spending on travel products is 2402, this value will be used in comparison against the train and test sets mean value of the dependent variable.

## Q-Q plot

```
travel_data %>% plot_qq()
```

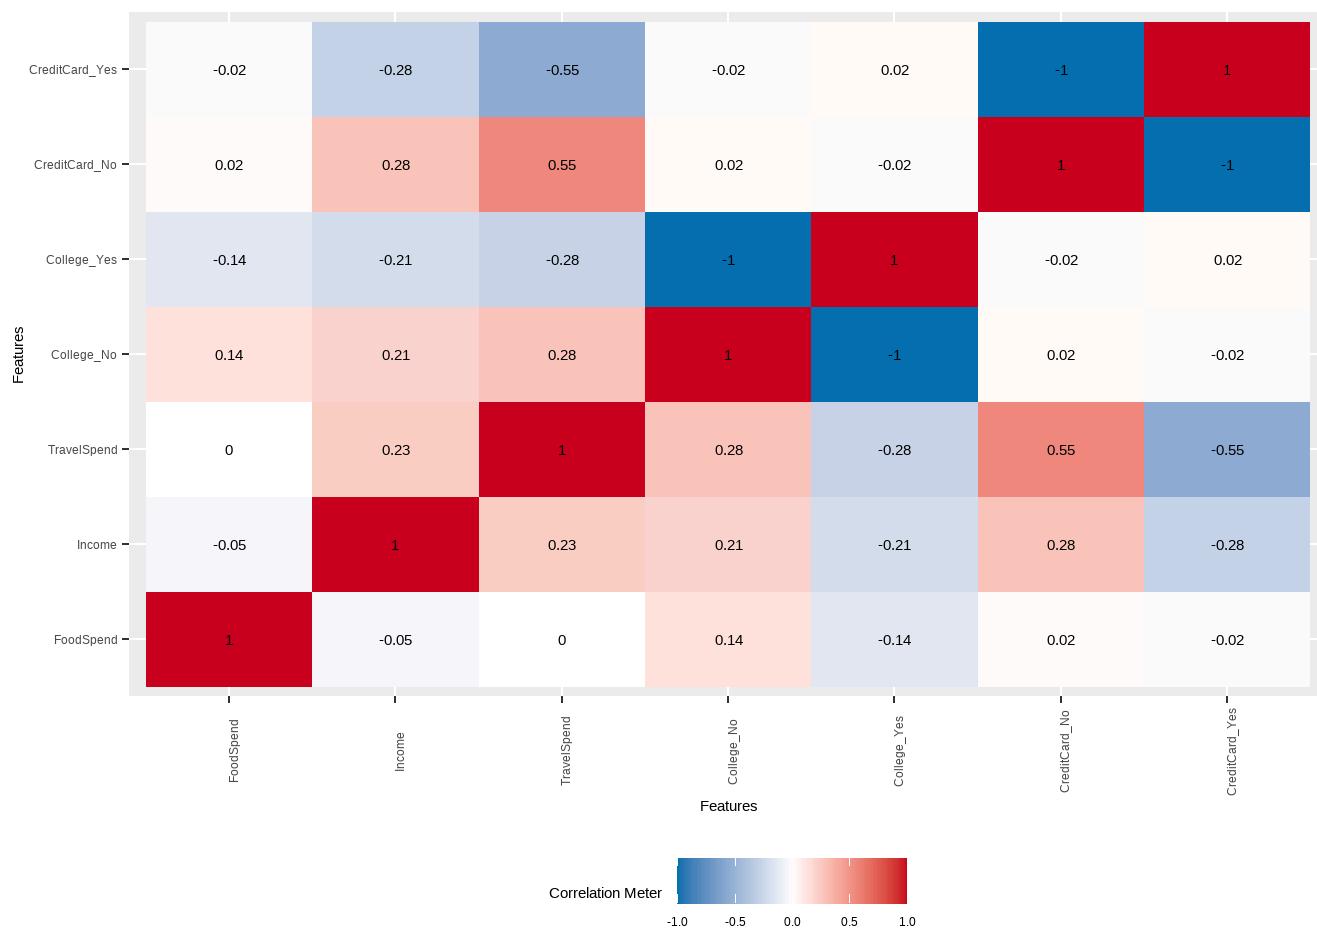


Comments on Q-Q plots:

The Q-Q plots of the numeric variables show that points follow closely along the diagonal reference line, with some deviations at the tail. This further supports that the variables are approximately normally distributed with mild right skewness.

## Correlation Matrix

```
travel_data %>% plot_correlation()
```



### Comments on Correlation Matrix:

This correlation matrix shows some interesting relationships among the variables. It is interesting that not having a credit card would lead to higher travel spend and income then not having a credit card, this relationships seems counterintuitive similarly, although not necessarily counterintuitive but interesting to note that not going to college also shows higher spending and income then people that did go to college. Travel Spend, has weak positive correlation with travel spend, people that did not attend college, and people that do not have a credit card.

## Data Preparation

### Variable Transformation

```
## Turn character variables into factors
travel_data$College <- as.factor(travel_data$College)
travel_data$CreditCard <- as.factor(travel_data$CreditCard)
travel_data %>% str()
```

```
tibble [500 x 5] (S3: tbl_df/tbl/data.frame)
$ College      : Factor w/ 2 levels "No","Yes": 2 1 1 1 1 1 1 1 1 ...
$ CreditCard   : Factor w/ 2 levels "No","Yes": 2 2 2 2 2 1 1 2 1 1 ...
```

```
$ FoodSpend : num [1:500] 5472 9131 4451 5021 3408 ...
$ Income     : num [1:500] 49150 47806 46050 42600 52092 ...
$ TravelSpend: num [1:500] 827 864 1180 756 851 ...
```

## Partition

```
set.seed(1)
myindex <- createDataPartition(travel_data$TravelSpend, p = 0.7, list = FALSE)
trainset <- travel_data[myindex,]
testset <- travel_data[-myindex,]

cat("Mean of the depvar Trainset")
```

Mean of the depvar Trainset

```
mean(trainset$TravelSpend)
```

[1] 2395.475

```
cat("Mean of the depvar Testset")
```

Mean of the depvar Testset

```
mean(testset$TravelSpend)
```

[1] 2418.489

Comments on Data Partition:

The modeling dataset was partitioned into a 70/30 split, 70% of the data will be used to train the model and 30% of the dataset will be used to test the model. The means of travel spend for the train and test dataset were both close which means both have similar distributions of the dependent variable.

## Modeling (Regression Trees)

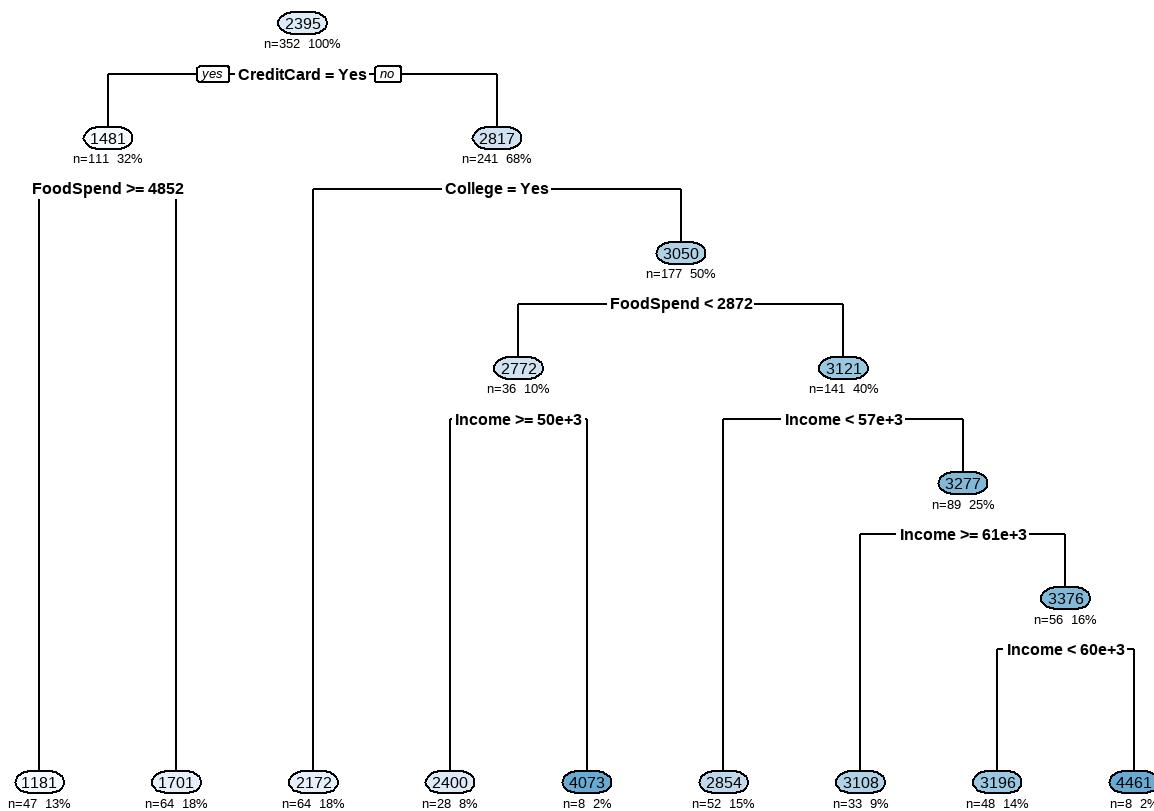
### Cross Validation

```
myctrl <- trainControl(method = "cv", number = 10)
```

### Default Tree (rpart)

```
## Default tree using cp = 0.01
set.seed(1)
```

```
tree_default <- rpart(TravelSpend~., data = trainset, method="anova")
rpart.plot(tree_default, type = 2, extra = 101, under = TRUE, fallen.leaves = TRUE)
```



```
## Find number of Leaf nodes in the default tree
cat("Number of Leaf Nodes in the Default Tree: ")
```

Number of Leaf Nodes in the Default Tree:

```
sum(tree_default$frame$var == "<leaf>")
```

[1] 9

Answer to Question 37 part A:

A) How many leaf nodes are in the default tree? What are the predictor variable and split value for the first split of the default regression tree?

- How many Leaf Nodes are in the default tree: There are 9 leaf nodes in the default tree.
- The predictor variable for the first split of the default regression tree is credit card. The split values are when credit card = yes the predicted mean Travel Spend is 1481, and if credit card = no then the predicted mean income is 2817.

## Full Grown Tree (rpart)

```
set.seed(1)
full_tree <- rpart(TravelSpend ~., data = trainset, control = rpart.control(cp = 0, minsplit =
    2, minbucket = 1, maxdepth = 30, xval = 10))

#xval = 10 is cross fold validation of 10, this is how to do it using rpart, used rpart to get
#correct values in the cp table

head(full_tree$cptable)
```

	CP	nsplit	rel error	xerror	xstd
1	0.30174408	0	1.0000000	1.0117564	0.06854470
2	0.08049551	1	0.6982559	0.7058791	0.05010272
3	0.01631113	2	0.6177604	0.6236386	0.04619739
4	0.01438172	3	0.6014493	0.6287099	0.04812099
5	0.01437510	8	0.5295407	0.6231581	0.04815957
6	0.01241362	10	0.5007905	0.6349426	0.05110058

## Where to Find Answers to Part B

```
min_error <- full_tree$cptable[which.min(full_tree$cptable[, "xerror"]), ]
min_error
```

	CP	nsplit	rel error	xerror	xstd
	0.01437510	8.0000000	0.52954066	0.62315809	0.04815957

Answer to Question 37 Part B:

B) Build a full-grown tree. Which cp value is associated with the lowest cross-validation error? How many splits are in the minimum-error tree? What is the minimum cross-validation error?

- CP value associated with the lowest cross-validation error: CP = 0.01437510.
- Number of splits that are in the minimum-error tree: 8.
- Minimum cross validation error: xerror = 0.62315809.

## Best Pruned Tree (rpart)

```
## To find the best pruned tree, you take the minimum error found in the minimum error tree and
# add it to the minimum errors associated xstd value. These two values added together
# makes the best pruned trees threshold. Now we look for the lowest number of splits
# within the xerror threshold.
```

```
cpt <- full_tree$cptable

imin      <- which.min(cpt[, "xerror"]) ## Grabs the row that the minimum error is on and defines
# it with a variable
xerr_min <- cpt[imin, "xerror"] ## Gets the xerror from row 5
xstd_min <- cpt[imin, "xstd"] ## gets the xstd from row 5
```

```
threshold <- xerr_min + xstd_min ## Creates the xerror threshold for best pruned tree.
threshold ## Lowest amount of splits while staying below an xerror value of 0.6713177
```

[1] 0.6713177

```
within_threshold <- which(cpt[, "xerror"] <= threshold)

best_pruned_index <- within_threshold[which.min(cpt[within_threshold, "nsplit"])]
```

best\_pruned\_cp <- cpt[best\_pruned\_index, "CP"]  
 best\_pruned\_xerror <- cpt[best\_pruned\_index, "xerror"]  
 best\_pruned\_nsplits <- cpt[best\_pruned\_index, "nsplit"]

```
cat("Best-Pruned Tree (1-SE Rule)\n")
```

Best-Pruned Tree (1-SE Rule)

```
cat("CP value:", best_pruned_cp, "\n")
```

CP value: 0.01631113

```
cat("Cross-validation error:", best_pruned_xerror, "\n")
```

Cross-validation error: 0.6236386

```
cat("Number of splits:", best_pruned_nsplits, "\n")
```

Number of splits: 2

Answer to Question 37 Part C:

c. Is there a simpler tree with a cross-validation error that is within one standard error of the minimum error? If there is, then which cp value is associated with the best-pruned tree?

Yes, there is a simpler tree with a cross-validation error within one standard error of the minimum error. The cp value associated with the best-pruned tree is 0.01631113.

## Best tree (rpart)

### Make Model

```
best_pruned_tree <- prune(full_tree, cp=0.01631113)
```

### Variable Importance

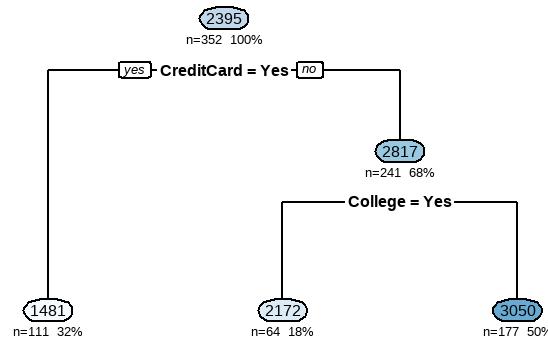
```
caret::varImp(best_pruned_tree)
```

Overall
College    0.20846187
CreditCard 0.30174408
FoodSpend 0.03065334
Income    0.06040484

Comments on Variable Importance:

Credit Card and College are the most important to the model. In the best pruned tree it can be expected that these variables will be predictor variables for splits.

```
rpart.plot(best_pruned_tree, type = 2, extra = 101, under = TRUE, fallen.leaves = TRUE)
```



```
cat("Number of Leaf Nodes in the Pruned Tree: ")
```

Number of Leaf Nodes in the Pruned Tree:

```
sum(best_pruned_tree$frame$var == "<leaf>")
```

[1] 3

Answer to Question 37 Part D:

D. Prune the full tree to the best-pruned tree or the minimum error tree if the answer to part c is "No."

Display the tree. How many leaf nodes are in the pruned tree?

There are three leaf nodes in the pruned tree.

Comments on Best Pruned Tree:

The Best pruned tree has Two Splits, The first split is whether or not the individual has a credit card, if the individual has a credit card then his annual travel spending is 1481. If the individual does not have a credit card then the annual travel spending is 2817. The second split is over whether an individual went to college or not, this split happens only if the person does not have a credit card. If the person did go to college the annual travel spending is 2172 whereas if the person did not go to college the annual travel spending is 3050.

## Evaluation (rpart)

### Evaluation Metrics

```

predict_bpt <- predict(best_pruned_tree, testset)
bpt_metrics <- round(forecast::accuracy(predict_bpt, testset$TravelSpend),2)

## For inline code
ME_value   <- bpt_metrics["Test set", "ME"]
RMSE_value <- bpt_metrics["Test set", "RMSE"]
MAE_value  <- bpt_metrics["Test set", "MAE"]
MPE_value  <- bpt_metrics["Test set", "MPE"]
MAPE_value <- bpt_metrics["Test set", "MAPE"]

rownames(bpt_metrics) <- "Rpart Regression Tree"
print(bpt_metrics)

```

ME	RMSE	MAE	MPE	MAPE
Rpart Regression Tree	-1.35	894.68	713.39	-14.75
				34.74

Answer to Question 37 Part E:

E) What are the ME, RMSE, MAE, MPE, and MAPE of the pruned tree on the validation data?

- **ME Value:** 'r ME\_value'.
- **RMSE Value:** 'r RMSE\_value'.
- **MAE Value:** 'r MAE\_value'.
- **MPE Value:** 'r MPE\_value' %.
- **MAPE Value:** 'r MAPE\_value' %.

Measures of Accuracy:

- **RMSE (Root Mean Square Error = 894.68)** and **MAE (Mean Absolute Error = 713.39)** show that the average prediction error of the pruned tree is roughly between **\$700–\$900** units from the actual value.

- **MAPE (Mean Absolute Percentage Error = 34.74%)** indicates that on average, predictions are off by about **35%** from the true travel spend values.
- Overall, the model demonstrates **moderate predictive accuracy** — it captures general trends but still has some variance in prediction magnitudes.

Measures of Bias:

- **ME (Mean Error = -1.35)** is close to zero, suggesting that the model is **unbiased** — it does not consistently over- or under-predict.
- **MPE (Mean Percentage Error = -14.75%)** is slightly negative, meaning the model **tends to slightly under-predict** actual values on average.

## Score (rpart)

```
score_rpart <- predict(best_pruned_tree, travel_score)
rpart_pip <- cbind(travel_score, score_rpart)
rpart_pip
```

	College	CreditCard	FoodSpend	Income	score_rpart
1	No	Yes	2892.90	65982	1480.546
2	Yes	No	6017.66	53986	2172.309

## Best Tree (caret)

```
set.seed(1)
best_tree <- train(TravelSpend ~.,
                    data = trainset,
                    method = "rpart",
                    trControl = myctrl,
                    tuneLength = 25,
                    metric = "RMSE",
                    control = rpart::rpart.control(minsplit = 2, minbucket = 1, cp = 0))

best_tree$resample %>%
  arrange(RMSE)
```

	RMSE	Rsquared	MAE	Resample
1	675.5641	0.6209010	513.9024	Fold08
2	783.7385	0.4410429	616.2088	Fold05
3	797.4464	0.5025473	562.8668	Fold10
4	864.7713	0.3207340	652.3549	Fold04
5	869.5082	0.4124079	691.6714	Fold02
6	908.2189	0.3394698	703.6228	Fold06
7	918.3013	0.4746556	688.2106	Fold09
8	973.0953	0.2522752	790.9866	Fold01

```
9 1015.4360 0.3027039 797.4441 Fold07
10 1025.5792 0.3539480 830.3901 Fold03
```

```
best_tree$finalModel$cptable
```

	CP	nsplit	rel error
1	0.30174408	0	1.0000000
2	0.08049551	1	0.6982559
3	0.01631113	2	0.6177604
4	0.01438172	3	0.6014493
5	0.01437510	8	0.5295407
6	0.01257267	10	0.5007905

## CP Values

```
best_tree$results
```

	cp	RMSE	Rsquared	MAE	RMSESD	RsquaredSD	MAESD
1	0.00000000	979.8144	0.3940203	609.5450	143.29946	0.12610666	109.18476
2	0.01257267	883.1659	0.4020686	684.7659	109.83998	0.11057390	102.95901
3	0.02514534	890.0210	0.3830431	717.3401	119.41866	0.13487652	105.91385
4	0.03771801	884.5279	0.3908443	715.2211	121.28057	0.13777580	105.54171
5	0.05029068	884.5279	0.3908443	715.2211	121.28057	0.13777580	105.54171
6	0.06286335	884.5279	0.3908443	715.2211	121.28057	0.13777580	105.54171
7	0.07543602	924.0716	0.3313033	761.2885	98.93832	0.12190218	89.51029
8	0.08800869	943.3660	0.3066676	778.3473	114.72822	0.11965967	102.89155
9	0.10058136	939.4593	0.3102737	783.7894	116.88836	0.11715062	100.17559
10	0.11315403	939.4593	0.3102737	783.7894	116.88836	0.11715062	100.17559
11	0.12572670	939.4593	0.3102737	783.7894	116.88836	0.11715062	100.17559
12	0.13829937	939.4593	0.3102737	783.7894	116.88836	0.11715062	100.17559
13	0.15087204	939.4593	0.3102737	783.7894	116.88836	0.11715062	100.17559
14	0.16344471	939.4593	0.3102737	783.7894	116.88836	0.11715062	100.17559
15	0.17601738	939.4593	0.3102737	783.7894	116.88836	0.11715062	100.17559
16	0.18859005	939.4593	0.3102737	783.7894	116.88836	0.11715062	100.17559
17	0.20116272	939.4593	0.3102737	783.7894	116.88836	0.11715062	100.17559
18	0.21373539	939.4593	0.3102737	783.7894	116.88836	0.11715062	100.17559
19	0.22630806	939.4593	0.3102737	783.7894	116.88836	0.11715062	100.17559
20	0.23888073	939.4593	0.3102737	783.7894	116.88836	0.11715062	100.17559
21	0.25145340	939.4593	0.3102737	783.7894	116.88836	0.11715062	100.17559
22	0.26402607	939.4593	0.3102737	783.7894	116.88836	0.11715062	100.17559
23	0.27659874	939.4593	0.3102737	783.7894	116.88836	0.11715062	100.17559
24	0.28917141	995.3909	0.2756474	830.0653	121.22026	0.10387564	102.28537
25	0.30174408	1042.4628	0.2371829	870.9928	89.14339	0.08918501	79.12694

```
cat("\nBest Tuned cp Value: ", best_tree$bestTune$cp)
```

Best Tuned cp Value: 0.01257267

## Variable Importance

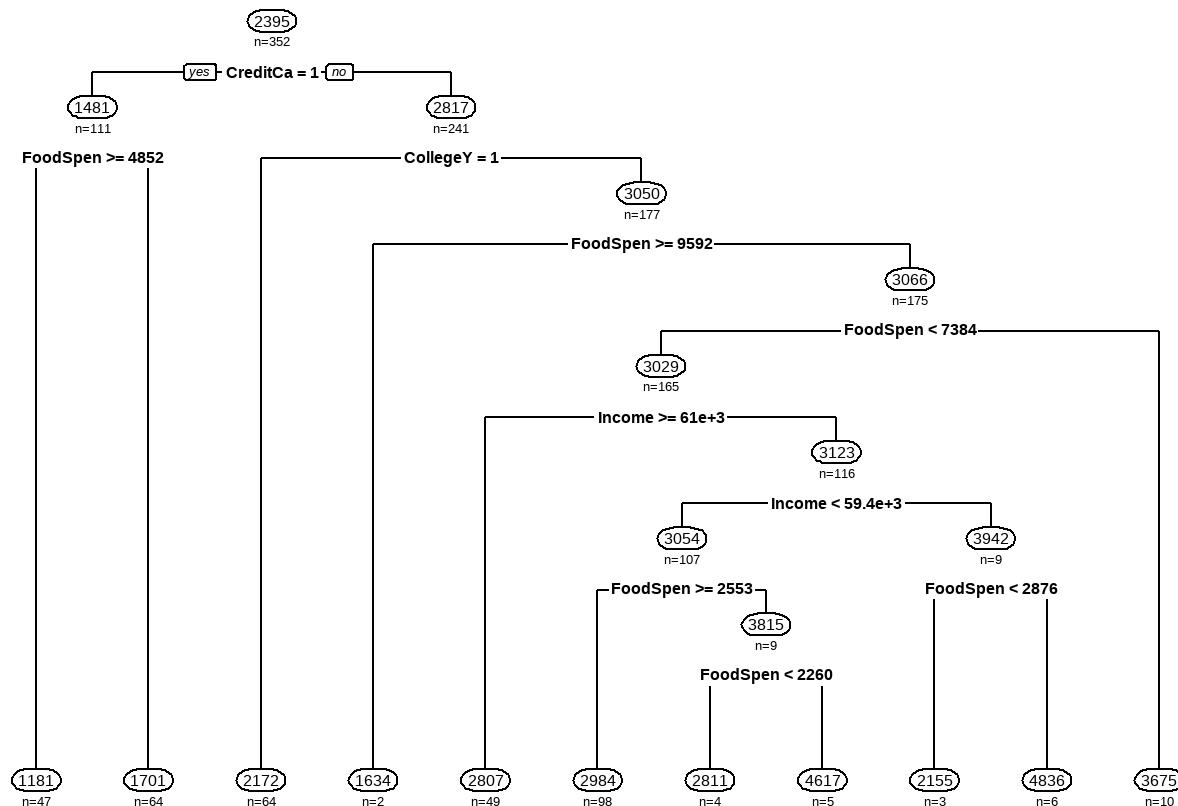
```
print(caret::varImp(best_tree))
```

rpart variable importance

Overall	
FoodSpend	100.000
Income	32.062
CreditCardYes	5.226
CollegeYes	0.000

## Show Best Pruned Tree

```
prp(best_tree$finalModel, type = 2, extra = 1, under = TRUE, digits = 3, fallen.leaves = TRUE)
```



## Number of Leaves in Best Tree (carert)

```
cat("Number of Leaf Nodes in the Best Tree: ")
```

Number of Leaf Nodes in the Best Tree:

```
sum(best_tree$finalModel$frame$var == "<leaf>")
```

[1] 11

## Evaluation (Caret Regression Tree)

```
predict_bp <- predict(best_tree, testset)
validation_metrics_tree <- round(forecast::accuracy(predict_bp, testset$TravelSpend), 2)
rownames(validation_metrics_tree) <- "Regression Tree"
print(validation_metrics_tree)
```

	ME	RMSE	MAE	MPE	MAPE
Regression Tree	-23.33	941.16	700.45	-14.27	32.47

## Scoring (Caret Regresssion Tree)

```
score_predict <- predict(best_tree, travel_score)
pip <- cbind(travel_score, score_predict)
knitr::kable(pip[order(pip$score_predict) , ])
```

College	CreditCard	FoodSpend	Income	score_predict
No	Yes	2892.90	65982	1700.855
Yes	No	6017.66	53986	2172.309

## Data Partition (Random Forest)

```
p <- ncol(trainset)-1

mtry_center <- max(1, round(p/3))
mtry_grid <- data.frame(mtry = sort(unique(pmax(1, c(mtry_center-1, mtry_center, mtry_center+1,
2, p)))))
```

## Modeling (Random Forest)

```
set.seed(1)
rf_model <- train(TravelSpend ~ .,
                   data = trainset,
                   method = "rf",
                   trControl = myctrl,
                   tuneGrid = mtry_grid,
```

```
ntree = 1000,
importance = TRUE)
```

## Resample

```
rf_model
```

Random Forest

352 samples  
4 predictor

No pre-processing  
Resampling: Cross-Validated (10 fold)  
Summary of sample sizes: 318, 317, 316, 316, 316, 318, ...  
Resampling results across tuning parameters:

mtry	RMSE	Rsquared	MAE
1	873.3978	0.4449263	715.9041
2	820.1100	0.4795871	644.1123
4	826.4942	0.4770504	607.5179

RMSE was used to select the optimal model using the smallest value.  
The final value used for the model was mtry = 2.

## Variable Importance

```
caret::varImp(rf_model)
```

rf variable importance

	Overall
CreditCardYes	100.000
CollegeYes	33.112
FoodSpend	3.641
Income	0.000

## Partial Effects of Predictors

```
## Predictors College, CreditCard, FoodSpend, Income

train_x <- subset(trainset, select = -TravelSpend)

pd_FoodSpend <- partial(object = rf_model,
                           pred.var = "FoodSpend",
```

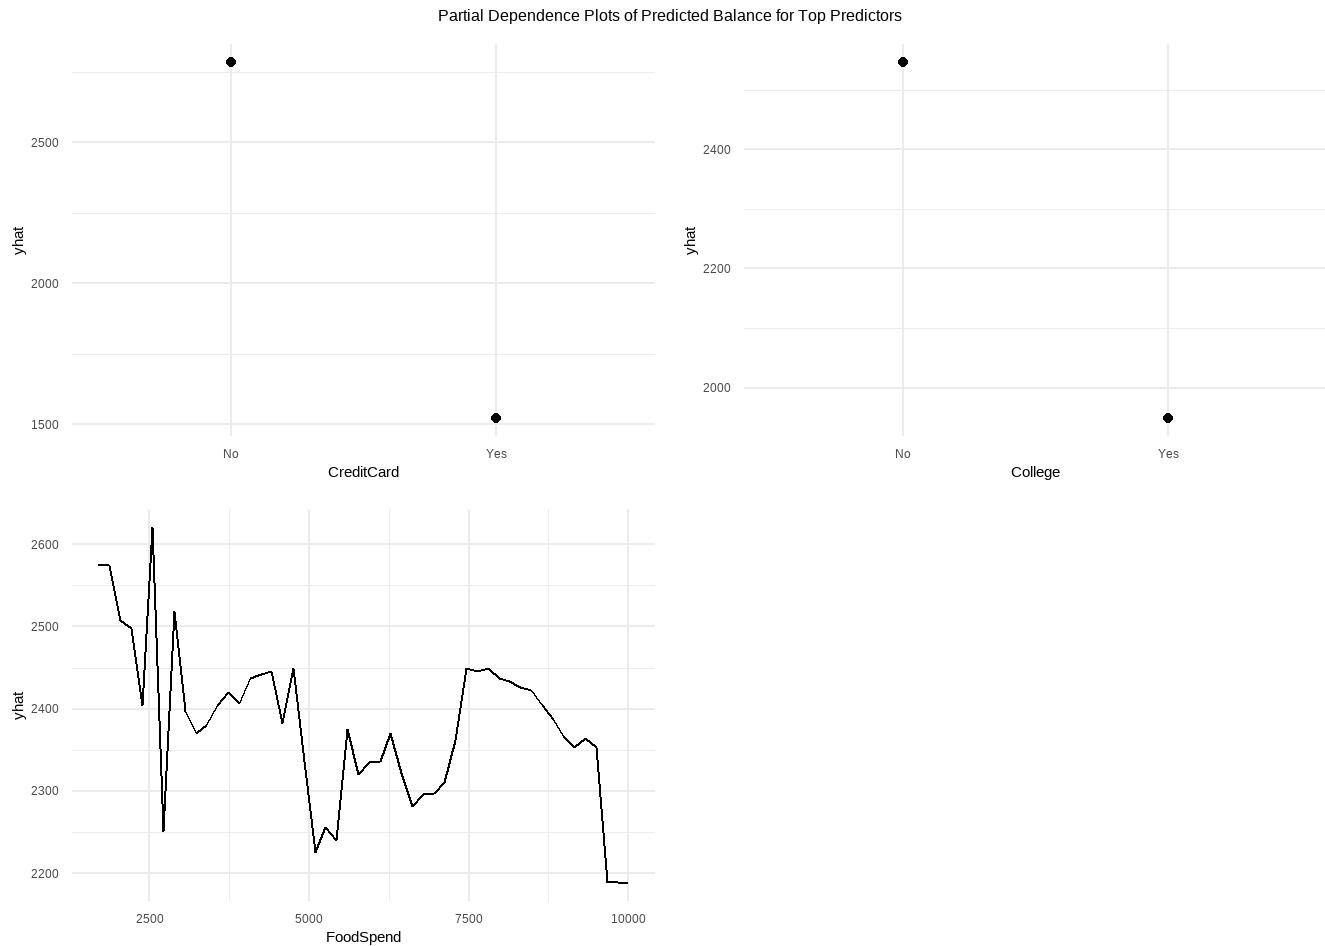
```

train = train_x,
grid.resolution = 50)
pd_ccd <- partial(object = rf_model,
pred.var = "CreditCard",
train = train_x,
grid.resolution = 50)
pd_college <- partial(object = rf_model,
pred.var = "College",
train = train_x,
grid.resolution = 50)

pd_fs <- autoplot(pd_FoodSpend) + theme_minimal()
pd_ccd <- autoplot(pd_ccd) + theme_minimal()
pd_col <- autoplot(pd_college) + theme_minimal()

grid.arrange(pd_ccd, pd_col, pd_fs, nrow = 2, ncol = 2, top = "Partial Dependence Plots of
Predicted Balance for Top Predictors")

```



## Evaluation (Random Forest)

```

predicted_rf <- predict(rf_model, testset)
test_rf <- round(forecast::accuracy(predicted_rf, testset$TravelSpend), 2)

```

```
rownames(test_rf) <- "Random Forest"
test_rf
```

	ME	RMSE	MAE	MPE	MAPE
Random Forest	-18.57	829.31	629.59	-12.79	29.32

## Compare Random Forest With Regression Tree

```
rbind(validation_metrics_tree, test_rf, bpt_metrics)
```

	ME	RMSE	MAE	MPE	MAPE
Regression Tree	-23.33	941.16	700.45	-14.27	32.47
Random Forest	-18.57	829.31	629.59	-12.79	29.32
Rpart Regression Tree	-1.35	894.68	713.39	-14.75	34.74

Comments Random Forest vs Regression Tree:

The Random Forest performs better than the Regression tree. The Random forest is more accurate than the regression tree as its RMSE, MAE, and MAPE are lower than the regression trees. The Random Forest also has minimal bias compared to the regression tree as the Random Forests ME and MPE are lower than the Regression trees ME and MPE.

## Score (Random Forest)

```
score_rf <- predict(rf_model, travel_score)
rf_pip <- cbind(travel_score, score_rf, score_predict, score_rpart)
rf_pip
```

	College	CreditCard	FoodSpend	Income	score_rf	score_predict	score_rpart
1	No	Yes	2892.90	65982	1816.319	1700.855	1480.546
2	Yes	No	6017.66	53986	1835.589	2172.309	2172.309

Comments on Comparing Random Forests to Regression Tree:

The Random Forest model's superior accuracy and low bias found in the key metrics suggest its travel spend predictions are more reliable than those of the regression tree. It is important to note that these predictions are not perfect and due to the Random Forests ME and MPE tend to slightly underpredict the true value.