

# Regression and Exponential Smoothing Time-Series Forecasting

## Lab 12 Forecasting

AUTHOR

Aaron Younger

PUBLISHED

November 21, 2025

## Business Understanding

In this analysis, non-seasonally adjusted sales time series data will be used to generate forecasts. The goal is to forecast August through December as accurately as possible.

## Data Understanding

### Correct R-Version

```
options(scipen = 999)
suppressWarnings(RNGversion("3.5.3"))
```

## Libraries

```
library(readxl)
library(DataExplorer)
library(tidyverse)
library(dplyr)
library(ggplot2)
library(strucchange)
library(forecast)
library(lubridate)
```

## Import Dataset

```
library(readxl)
my_data <- read_excel("MRTSSM451USN.xlsx",
    sheet = "Monthly")
View(my_data)

my_data %>% head(1) ## Beginning date 1992-01-01
```

```
# A tibble: 1 × 2
observation_date    MRTSSM451USN
<dttm>                  <dbl>
1 1992-01-01 00:00:00      2982
```

```
my_data %>% tail(1) ## End date 2025-07-01
```

```
# A tibble: 1 × 2
observation_date    MRTSSM451USN
<dttm>                  <dbl>
1 2025-07-01 00:00:00      8010
```

```
my_data <- my_data %>%
  rename(sales = MRTSSM451USN)
```

Comments on Import Dataset:

After importing the forecast data that will be used for analysis, I checked the beginning and end date of the dataset to know the range of dates. The beginning date in this dataset is January 1, 1992 and the end date in this dataset is July 1, 2025.


## Create New Forecasting Dataset

```
my_data <- my_data %>%
  arrange(observation_date)

start_year  <- year(min(my_data$observation_date))
start_month <- month(min(my_data$observation_date))

forecast_data <- ts(
  my_data$sales,
  start     = c(start_year, start_month),
  frequency = 12
)

plot.ts(forecast_data)
```



Comments on Creation of time series dataset:


Based on the graph of sales throughout the dates, it is clear there is both an upward trend and yearly seasonality during the end of each year. There is also a significant drop in sales around 2019-2021. This is due to COVID-19 and needs to be addressed during analysis for a potential structure break.

## Quantitatively show the affects of seasonality on sales

```
monthly_avg <- my_data %>%
  mutate(month = month(observation_date, label = TRUE, abbr = TRUE)) %>%
  group_by(month) %>%
  summarise(avg_sales = mean(sales))

ggplot(data = monthly_avg, aes(x=month, y = avg_sales)) +
  geom_col(fill = "blue") +
  geom_text(aes(label = round(avg_sales,0)),
            vjust = -0.5) +
  labs(
    title = "Average Sales by Month",
    subtitle = "Investigate Seasonality",
    x = "Month",
    y = "Average Sales"
```

```
) +  
theme_minimal()
```




Comments on Seasonality:

It is apparent from the bar chart that December has by far the highest average sales. This is most likely due to the holiday season.

## Bai-Perron, test for Structure Breaks

```
sb <- strucchange::breakpoints(forecast_data ~ 1)  
plot(sb) ## Figure out what this graph reads
```

## BIC and Residual Sum of Squares



```
strucchange::breakdates(sb)
```

```
[1] 1997.750 2004.750 2020.333
```

```
confint(sb)
```

Confidence intervals for breakpoints  
of optimal 4-segment partition:

Call:

```
confint.breakpointsfull(object = sb)
```

Breakpoints at observation number:

2.5 % breakpoints 97.5 %

	59	70	80
1	136	154	172
2	334	341	350


Corresponding to breakdates:

2.5 % breakpoints 97.5 %

```
1 1996(11) 1997(10) 1998(8)
```

2 2003(4)	2004(10)	2006(4)
3 2019(10)	2020(5)	2021(2)

```
plot(forecast_data, main = "Sales with Detected Structure Breaks")
lines(sb, col = "red", lwd = 2)
```



#### Comments on Bai-Perron Test:

The Bai-Perron test is a statistical test that shows breakpoints in the forecasting data. Breakpoints in forecast data refer to significant changes or discontinuities that occur in the data, which can affect the accuracy and reliability of forecasts. To look for breakpoints I first graphed a BIC and RSS graph. RSS, the blue line, measures model error, whereas the BIC, the black line, penalizes model complexity. The best model is the one with the lowest BIC, BIC will drop as breakpoints are added but will rise again once too many are added. This BIC and RSS graph show the optimal number of breakpoints is 3. The optimal number of breakpoints is further supported by the confidence interval table for structural break dates. The estimated breakdates for the three detected structural changes fall within the following ranges: November 1996 to August 1998, April 2003 to April 2006, and October 2019 to February 2021. These breakpoints are visually displayed on the plotted forecast series, where they appear as red vertical lines intersecting the time series curve.

#### Chow Test

```
chow_test <- sctest(forecast_data ~ 1, type = "Chow", point = 334)
chow_test
```

Chow test

```
data: forecast_data ~ 1
F = 97.836, p-value < 0.0000000000000022
```

Comments on Chow Test:

The Chow Test tests for whether a given date represents a statistically significant structural break in the time series data. Since I want the time series to start as early as possible while still being structurally consistent I will use the Chow test on October, 2019. The results of the chow test showed that October of 2019 is significant due to the extremely small p-value. What this means is for modeling my new start date will be October 1, 2019.

## Data Preparation: Regression

### Create Dummy Variable for Covid Era

```
covid_start <- c(2020, 3)    # March 2020
covid_end   <- c(2021, 12)   # December 2021

# Create the dummy on the same time index as forecast_data
covid_dummy <- ifelse(
  time(forecast_data) >= covid_start[1] + (covid_start[2] - 1) / 12 &
  time(forecast_data) <= covid_end[1] + (covid_end[2] - 1) / 12,
  1, 0
)

covid_dummy <- ts(covid_dummy,
                  start = start(forecast_data),
                  frequency = frequency(forecast_data))
```

Comments on Creating Dummy Variable for Covid Era:

A binary dummy variable was constructed to represent the COVID-19 era. Observations occurring during the defined COVID period were assigned a value of 1, and all other observations were assigned a 0. Including this variable allows the model to test whether the COVID-era period had a statistically significant effect by comparing model performance with and without the dummy term. The COVID-19 era is defined as March 2020 - December 2021 since this data range captures the full shock and full recovery from COVID-19.

### Partition, forecast data

```
trainset <- window(forecast_data, start = c(2019,10), end =c(2024,6))
testset <- window(forecast_data, start = c(2024,7), end = c(2025,7))

cat("trainset has", length(trainset), "observation.")
```


trainset has 57 observation.

```
cat("\ntestset has", length(testset), "observation.")
```

testset has 13 observation.

```
plot_set <- bind_rows(data.frame(Time = time(trainset), sales = as.numeric(trainset), set =
  "Train"),
  data.frame(Time = time(testset), sales = as.numeric(testset), set =
  "Test"))
ggplot(plot_set, aes(Time, sales, color = set)) +
  geom_line(linewidth = 1) +
  geom_vline(xintercept = 2024.5, linetype = 2) +
  labs(title = "Sales Data by Month: Train vs. Test Split",
  x = "Time", y = "Sales") +
  theme_minimal()
```

Sales Data by Month: Train vs. Test Split



## Comments on Partitioning the Forecast Data:

A training and test dataset was made for modeling. The training dataset ranged from October 2019 - June 2024. The test dataset was exactly one year, it ranged from July 2024 to the end of the time series dataset July 2025.

## Partition, covid binary

```
covid_train <- window(covid_dummy, start = c(2019, 10), end = c(2024, 6))
```

## Modeling: Regression

### Cross Fold Validation: Regression

```
t_results <- data.frame() # store all folds + models here

# FOLD 1:

# y (sales)
cv_train1_y <- window(trainset, end = c(2021,12))
cv_val1_y   <- window(trainset, start = c(2022,1), end = c(2022,12))

# covid dummy (aligned)
cv_train1_c <- window(covid_train, end = c(2021,12))
cv_val1_c   <- window(covid_train, start = c(2022,1), end = c(2022,12))

## Model 1: trend
model1_M1    <- tslm(cv_train1_y ~ trend)
forecast1_M1 <- forecast(model1_M1, h = length(cv_val1_y))
acc1_M1      <- forecast::accuracy(forecast1_M1, cv_val1_y)

t_results <- rbind(t_results,
  data.frame(Fold = 1,
    Model = "trend",
    RMSE  = acc1_M1[2,"RMSE"],
    MAPE  = acc1_M1[2,"MAPE"]))
)

## Model 2: trend + season
model1_M2    <- tslm(cv_train1_y ~ trend + season)
forecast1_M2 <- forecast(model1_M2, h = length(cv_val1_y))
acc1_M2      <- forecast::accuracy(forecast1_M2, cv_val1_y)

t_results <- rbind(t_results,
  data.frame(Fold = 1,
    Model = "trend + season",
```

```

RMSE  = acc1_M2[2,"RMSE"],
MAPE  = acc1_M2[2,"MAPE"])
)

## Model 3: trend + season + covid
model1_M3    <- tslm(cv_train1_y ~ trend + season + cv_train1_c)
forecast1_M3 <- forecast(
  model1_M3,
  h = length(cv_val1_y),
  newdata = data.frame(cv_train1_c = cv_val1_c)
)
acc1_M3      <- forecast::accuracy(forecast1_M3, cv_val1_y)

t_results <- rbind(t_results,
  data.frame(Fold = 1,
    Model = "trend + season + covid",
    RMSE  = acc1_M3[2,"RMSE"],
    MAPE  = acc1_M3[2,"MAPE"]))
)

# FOLD 2:

cv_train2_y <- window(trainset,   end   = c(2022,12))
cv_val2_y   <- window(trainset,   start = c(2023,1), end = c(2023,12))

cv_train2_c <- window(covid_train, end   = c(2022,12))
cv_val2_c   <- window(covid_train, start = c(2023,1), end = c(2023,12))

## Model 1: trend
model2_M1    <- tslm(cv_train2_y ~ trend)
forecast2_M1 <- forecast(model2_M1, h = length(cv_val2_y))
acc2_M1      <- forecast::accuracy(forecast2_M1, cv_val2_y)

t_results <- rbind(t_results,
  data.frame(Fold = 2,
    Model = "trend",
    RMSE  = acc2_M1[2,"RMSE"],
    MAPE  = acc2_M1[2,"MAPE"]))
)

## Model 2: trend + season
model2_M2    <- tslm(cv_train2_y ~ trend + season)
forecast2_M2 <- forecast(model2_M2, h = length(cv_val2_y))
acc2_M2      <- forecast::accuracy(forecast2_M2, cv_val2_y)

t_results <- rbind(t_results,
  data.frame(Fold = 2,
    Model = "trend + season",
    RMSE  = acc2_M2[2,"RMSE"],
    MAPE  = acc2_M2[2,"MAPE"]))
)

```

```

MAPE  = acc2_M2[2,"MAPE"])
)

## Model 3: trend + season + covid
model2_M3    <- tslm(cv_train2_y ~ trend + season + cv_train2_c)
forecast2_M3 <- forecast(
  model2_M3,
  h = length(cv_val2_y),
  newdata = data.frame(cv_train2_c = cv_val2_c)
)
acc2_M3      <- forecast::accuracy(forecast2_M3, cv_val2_y)

t_results <- rbind(t_results,
  data.frame(Fold = 2,
    Model = "trend + season + covid",
    RMSE  = acc2_M3[2,"RMSE"],
    MAPE  = acc2_M3[2,"MAPE"]))
)

# FOLD 3:

cv_train3_y <- window(trainset,   end   = c(2023,12))
cv_val3_y   <- window(trainset,   start = c(2024,1), end = c(2024,6))

cv_train3_c <- window(covid_train, end   = c(2023,12))
cv_val3_c   <- window(covid_train, start = c(2024,1), end = c(2024,6))

## Model 1: trend
model3_M1    <- tslm(cv_train3_y ~ trend)
forecast3_M1 <- forecast(model3_M1, h = length(cv_val3_y))
acc3_M1      <- forecast::accuracy(forecast3_M1, cv_val3_y)

t_results <- rbind(t_results,
  data.frame(Fold = 3,
    Model = "trend",
    RMSE  = acc3_M1[2,"RMSE"],
    MAPE  = acc3_M1[2,"MAPE"]))
)

## Model 2: trend + season
model3_M2    <- tslm(cv_train3_y ~ trend + season)
forecast3_M2 <- forecast(model3_M2, h = length(cv_val3_y))
acc3_M2      <- forecast::accuracy(forecast3_M2, cv_val3_y)

t_results <- rbind(t_results,
  data.frame(Fold = 3,
    Model = "trend + season",
    RMSE  = acc3_M2[2,"RMSE"],
    MAPE  = acc3_M2[2,"MAPE"]))
)

```

```

## Model 3: trend + season + covid
model3_M3      <- tslm(cv_train3_y ~ trend + season + cv_train3_c)
forecast3_M3 <- forecast(
  model3_M3,
  h = length(cv_val3_y),
  newdata = data.frame(cv_train3_c = cv_val3_c)
)
acc3_M3       <- forecast::accuracy(forecast3_M3, cv_val3_y)

t_results <- rbind(t_results,
  data.frame(Fold = 3,
    Model = "trend + season + covid",
    RMSE  = acc3_M3[2,"RMSE"],
    MAPE  = acc3_M3[2,"MAPE"]))
)

t_results

```

	Fold	Model	RMSE	MAPE
1	1	trend	2020.458	24.94214
2	1	trend + season	1329.589	15.56136
3	1	trend + season + covid	2451.921	29.64513
4	2	trend	1804.887	22.12157
5	2	trend + season	1190.354	14.07086
6	2	trend + season + covid	1143.092	13.40971
7	3	trend	1950.668	26.63008
8	3	trend + season	1049.182	13.65372
9	3	trend + season + covid	1132.616	14.98762

Comments on Cross Fold Validation:

Three different models were tested over three different time periods. Each models RMSE and MAPE was collected from each cross validation fold. The RMSE and MAPE will be evaluated and whichever model has the lowest average RMSE and MAPE from the three fold cross validation will be the model used to forecast.

## Evaluation: Regression

```

# Average CV error per model (this is what you report)
cv_summary <- t_results %>%
  group_by(Model) %>%
  summarise(
    mean_RMSE = mean(RMSE),
    mean_MAPE = mean(MAPE),
    .groups = "drop"
  )

cv_summary

```

```
# A tibble: 3 × 3
  Model          mean_RMSE  mean_MAPE
  <chr>        <dbl>      <dbl>
1 trend         1925.       24.6
2 trend + season 1190.     14.4
3 trend + season + covid 1576.     19.3
```

Comments on Evaluation:

The model with the lowest mean RMSE and MAPE is the model that includes just trend and season. This Model will be used to forecast the sales of August - December 2025.

## Deployment: Regression

### Forecasting 2025-08-01 - 2025-12-01

```
newdata2025 <- window(forecast_data, start = c(2019,5), end = c(2025,7))
new_covid <- window(covid_dummy, start = c(2019,5), end = c(2025,7))

final_model1 <- tslm(newdata2025 ~ trend + season)
summary(final_model1)
```

Call:

```
tslm(formula = newdata2025 ~ trend + season)
```

Residuals:

Min	1Q	Median	3Q	Max
-2841.58	-436.25	61.15	468.95	1643.19

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	5631.900	347.525	16.206	< 0.0000000000000002 ***
trend	23.853	4.063	5.871	0.00000018362276058 ***
season2	-449.520	437.439	-1.028	0.30812
season3	849.961	437.496	1.943	0.05659 .
season4	187.441	437.590	0.428	0.66988
season5	833.396	421.587	1.977	0.05252 .
season6	1274.258	421.529	3.023	0.00364 **
season7	1129.405	421.509	2.679	0.00943 **
season8	1876.432	437.892	4.285	0.00006478977553863 ***
season9	885.079	437.722	2.022	0.04750 *
season10	794.559	437.590	1.816	0.07424 .
season11	1985.706	437.496	4.539	0.00002658858924580 ***
season12	4495.853	437.439	10.278	0.0000000000000512 ***
---				
Signif. codes:	0 ****	0.001 **	0.01 *'	0.05 '.'
	0.1	'	'	1

Residual standard error: 757.6 on 62 degrees of freedom  
 Multiple R-squared: 0.7752, Adjusted R-squared: 0.7317  
 F-statistic: 17.82 on 12 and 62 DF, p-value: 0.000000000000000884


```
future_covid <- data.frame(new_covid = rep(0, 5))

forecast <- forecast(final_model1, h = 5, newdata = future_covid)
forecast
```

	Point Forecast	Lo 80	Hi 80	Lo 95	Hi 95
Aug 2025	9321.159	8238.320	10403.999	7650.140	10992.179
Sep 2025	8353.659	7270.820	9436.499	6682.640	10024.679
Oct 2025	8286.993	7204.153	9369.832	6615.973	9958.012
Nov 2025	9501.993	8419.153	10584.832	7830.973	11173.012
Dec 2025	12035.993	10953.153	13118.832	10364.973	13707.012

```
autoplot(forecast) +
  labs(title = "Forecast vs. Historical Data") +
  theme_minimal()
```

### Forecast vs. Historical Data



Comments on Forecasting Using Regression:

Overall, the final model accounts for 73 percent variability within the dataset. The final regression model predicts that August will have a sales value of 9,321, September will have a sales value of 8,354, October will have a sales value of 8,287, November will have a sales value of 9,502, December will have a sales value of 12,036. These estimates are reasonable as they follow a similar trend to the other Aug-Dec months in the dataset.


## Data Understanding: Smoothing

### Create data time period for Smoothing

```
my_data_recent <- my_data %>%
  filter(observation_date >= as.Date("2019-10-01")) %>%
  arrange(observation_date)

start_year1 <- year(min(my_data_recent$observation_date))
start_month1 <- month(min(my_data_recent$observation_date))

smoothing_data <- ts(my_data_recent$sales,
                      start = c(start_year1, start_month1), frequency = 12)
plot(smoothing_data)
```



### Comments on smoothing data:


For the smoothing dataset I will use the data that is found in the most recent structure break, cutting off all dates behind this. Including older dates that also have their own structure breaks will distort forecast accuracy of the smoothing models. The most recent structure break and the following data is most likely to produce the most accurate forecasts.

### Comments on trend and seasonality:

Based on the time-series graph, there is seasonality. This is due to the spike in sales during December. To confirm trend and or seasonality I will use another method that tests for seasonality and trend.

## STL Decomposition; test for both trend and seasonality

```
fit <- stl(smoothing_data, s.window = "periodic")
autoplot(fit)
```



### Comments on STL Decomposition Graphs:

Plotting STL produces four different graphs. The top graph labeled data, shows the raw time series data. The second graph labeled trend shows a strong trend upward from 2020 to mid-2022, most likely representing the shift out of COVID. The trend then flattens through 2022 to 2025, with there being a slight increase from 2024 to 2025. The third graph labeled seasonal shows a repeating yearly cycle of high sales.

The fourth graph remainder shows what the time series graph looks like when trend and seasonality is removed. Overall, it is safe to say that the period of time series data that smoothing will be done over does contain both trend and seasonality. For modeling this means Holt-Winters should be used.

## Modeling: Holt-Winters Smoothing

### Cross Validation - Holt-Winters Smoothing

```
s_results <- data.frame()

## FOLD 1: train <= 2021-12, validate 2022-01 .. 2022-12
s_train1 <- window(smoothing_data, end = c(2021, 12))
s_val1   <- window(smoothing_data, start = c(2022, 1), end = c(2022, 12))

fit1_AAA <- ets(s_train1, model = "AAA", restrict = FALSE)
fc1_AAA  <- forecast(fit1_AAA, h = length(s_val1))
acc1_AAA <- accuracy(fc1_AAA, s_val1)

s_results <- rbind(s_results,
  data.frame(Fold = 1, Model = "AAA",
             RMSE = acc1_AAA[2, "RMSE"],
             MAPE = acc1_AAA[2, "MAPE"]))
)

fit1_AAM <- ets(s_train1, model = "AAM", restrict = FALSE)
fc1_AAM  <- forecast(fit1_AAM, h = length(s_val1))
acc1_AAM <- accuracy(fc1_AAM, s_val1)

s_results <- rbind(s_results,
  data.frame(Fold = 1, Model = "AAM",
             RMSE = acc1_AAM[2, "RMSE"],
             MAPE = acc1_AAM[2, "MAPE"]))
)

## FOLD 2: train <= 2022-12, validate 2023
s_train2 <- window(smoothing_data, end = c(2022, 12))
s_val2   <- window(smoothing_data, start = c(2023, 1), end = c(2023, 12))

fit2_AAA <- ets(s_train2, model = "AAA", restrict = FALSE)
fc2_AAA  <- forecast(fit2_AAA, h = length(s_val2))
acc2_AAA <- accuracy(fc2_AAA, s_val2)

s_results <- rbind(s_results,
  data.frame(Fold = 2, Model = "AAA",
             RMSE = acc2_AAA[2, "RMSE"],
             MAPE = acc2_AAA[2, "MAPE"]))
)
```

```

fit2_AAM <- ets(s_train2, model = "AAM", restrict = FALSE)
fc2_AAM  <- forecast(fit2_AAM, h = length(s_val2))
acc2_AAM <- accuracy(fc2_AAM, s_val2)

s_results <- rbind(s_results,
  data.frame(Fold = 2, Model = "AAM",
             RMSE = acc2_AAM[2, "RMSE"],
             MAPE = acc2_AAM[2, "MAPE"]))
)

## FOLD 3: train <= 2023-12, validate 2024-01 .. 2024-06
s_train3 <- window(smoothing_data, end = c(2023, 12))
s_val3   <- window(smoothing_data, start = c(2024, 1), end = c(2024, 6))

fit3_AAA <- ets(s_train3, model = "AAA", restrict = FALSE)
fc3_AAA  <- forecast(fit3_AAA, h = length(s_val3))
acc3_AAA <- accuracy(fc3_AAA, s_val3)

s_results <- rbind(s_results,
  data.frame(Fold = 3, Model = "AAA",
             RMSE = acc3_AAA[2, "RMSE"],
             MAPE = acc3_AAA[2, "MAPE"]))
)

fit3_AAM <- ets(s_train3, model = "AAM", restrict = FALSE)
fc3_AAM  <- forecast(fit3_AAM, h = length(s_val3))
acc3_AAM <- accuracy(fc3_AAM, s_val3)

s_results <- rbind(s_results,
  data.frame(Fold = 3, Model = "AAM",
             RMSE = acc3_AAM[2, "RMSE"],
             MAPE = acc3_AAM[2, "MAPE"]))
)

s_results

```

	Fold	Model	RMSE	MAPE
1	1	AAA	1869.4543	21.501937
2	1	AAM	1752.9373	19.919131
3	2	AAA	1406.0059	16.153866
4	2	AAM	397.0064	3.860729
5	3	AAA	638.2105	7.069445
6	3	AAM	387.9680	4.073419

Comments on Cross Fold Validation:

Two different types of Holt-Winter models were tested. The AAA smoothing method is the additive version of the Holt-Winters seasonal smoothing and AAM is the multiplicative Holt-Winters smoothing. The AAA model is used when seasonality stays constant and AAM is used when the size of seasonal swing grows or shrinks as the time changes. Each smoothing models RMSE and MAPE was collected from each cross

validation fold. The RMSE and MAPE will be evaluated and whichever model has the lowest average RMSE and MAPE from cross validation will be the model used to forecast.

## Evaluation: Holt-Winters Smoothing

```
s_results_summary <- s_results %>%
  group_by(Model) %>%
  summarise(
    mean_RMSE = mean(RMSE),
    mean_MAPE = mean(MAPE),
    .groups = "drop"
  )
s_results_summary
```

```
# A tibble: 2 × 3
  Model mean_RMSE mean_MAPE
  <chr>     <dbl>     <dbl>
1 AAA      1305.     14.9
2 AAM      846.      9.28
```

Comments on Evaluation:

The Model with the lowest RMSE and MAPE is the AAM version of Holt-Winters. This model will be used to forecast the sales of August - December 2025.


## Deployment - Forecast sales for August - December 2025

```
sfinal <- ets(smoothing_data, model = "AAM", restrict = FALSE)
sf_fcst <- forecast::forecast(sfinal, h=5)
sf_fcst
```

	Point Forecast	Lo 80	Hi 80	Lo 95	Hi 95
Aug 2025	8544.612	7787.469	9320.988	7360.140	9706.227
Sep 2025	7575.036	6734.312	8409.806	6277.007	8839.335
Oct 2025	7512.327	6571.779	8468.210	6075.740	8935.495
Nov 2025	8869.347	7766.350	9962.485	7209.020	10546.128
Dec 2025	11392.514	10055.070	12779.738	9342.565	13461.565

```
plot(sf_fcst)
```

## Forecasts from ETS(A,Ad,M)



Comments on Forecasting using Smoothing:

The Final AAM Holt-Winters smoothing model predicts that August will have a sales value of 8545, September will have a sales value of 7,575, October will have a sales value of 7,512, November will have a sales value of 8,869, December will have a sales value of 11,393. These estimates are reasonable as they follow a similar trend to the other Aug-Dec months in the dataset.

## Which is the Better Model: Regression vs Holt-Winters Smoothing

```
model_compare <- data.frame(
  Model = c("Regression", "Holt-Winters AAM"),
  RMSE = c(mean(c(acc1_M2[2,"RMSE"], acc2_M2[2,"RMSE"], acc3_M2[2,"RMSE"])),
           mean(c(acc1_AAM[2,"RMSE"], acc2_AAM[2,"RMSE"], acc3_AAM[2,"RMSE"]))),
  MAPE = c(mean(c(acc1_M2[2,"MAPE"], acc2_M2[2,"MAPE"], acc3_M2[2,"MAPE"])),
           mean(c(acc1_AAM[2,"MAPE"], acc2_AAM[2,"MAPE"], acc3_AAM[2,"MAPE"])))
)
model_compare
```

	Model	RMSE	MAPE
1	Regression	1189.7087	14.428648

2 Holt-Winters AAM 845.9706 9.284426

#### Comments on Best Model:

Based off the best model for both Regression and Holt-Winters Smoothing, Holt-Winters AAM should be used to forecast and not the best regression model. Holt-Winters AAM is preferred over regression for forecasting August-December 2025 because it has a lower RMSE and MAPE. This means the Holt-Winters AAM has lower average error producing more accurate sales numbers.