

Lab 10: Regression Tress, Random Forest (Problem 45)

AUTHOR

Aaron Younger

Business Understanding

Merrick Stevens is a sports analyst working for ACE Sports Management, a sports agency that represents over 200 athletes. He is interested in understanding the relationship between an NBA player's salary and his physicality and performance statistics.

Data Understanding

R Version

```
options(scipen = 999)
suppressWarnings(RNGversion("3.5.3"))
```

Libraries

```
library(janitor)
library(DataExplorer)
library(readxl)
library(tidyverse)
library(ggplot2)
library(dplyr)
library(e1071)
library(dlookr)
library(caret)
library(rpart)
library(rpart.plot)
library(randomForest)
```

Import Modeling and Score Dataset

```
nba_data <- read_excel("jaggia_ba_2e_ch13_data.xlsx",
  sheet = "NBA_Data")
nba_data <- clean_names(nba_data)
View(nba_data)

nba_score <- read_excel("jaggia_ba_2e_ch13_data.xlsx",
```

```
sheet = "NBA_Score")
nba_score <- clean_names(nba_score)
```

Data Set Key:

- player_number: Represents the identification number assigned to each NBA player.
- salary (Dependent Variable) : Indicates the annual salary earned by the NBA player (in dollars).
- age: Shows the player's age in years.
- height: Represents the player's height (typically measured in inches).
- weight: Represents the player's weight (typically measured in pounds).
- games_played: The total number of games the player has participated in during the season.
- games_started: The total number of games in which the player was part of the starting lineup.
- minutes_per_game: The average number of minutes the player spends on the court per game.
- FG_made: Total number of field goals successfully made by the player.
- FG_attempted: Total number of field goal attempts taken by the player.
- FG_percent: Field goal shooting percentage ($FG_made \div FG_attempted \times 100$).
- 3P_made: Total number of three-point shots successfully made by the player.
- 3P_attempted: Total number of three-point shot attempts taken by the player.
- 3P_percent: Three-point shooting percentage ($3P_made \div 3P_attempted \times 100$).
- FT_made: Total number of free throws successfully made by the player.
- FT_attempted: Total number of free throw attempts taken by the player.
- FT_percent: Free throw shooting percentage ($FT_made \div FT_attempted \times 100$).
- offensive_rebounds: The total number of rebounds collected on the offensive side of the court.
- defensive_rebounds: The total number of rebounds collected on the defensive side of the court.
- assists: The number of times a player passes the ball leading directly to a made basket.
- blocks: The total number of opponent shots blocked by the player.
- steals: The total number of times the player takes the ball away from an opponent.
- personal_fouls: The total number of personal fouls committed by the player.
- turnovers: The total number of times the player loses possession of the ball to the opposing team.
- points: The total number of points scored by the player throughout the season.

Dataset Exploration

```
nba_data %>% head()
```

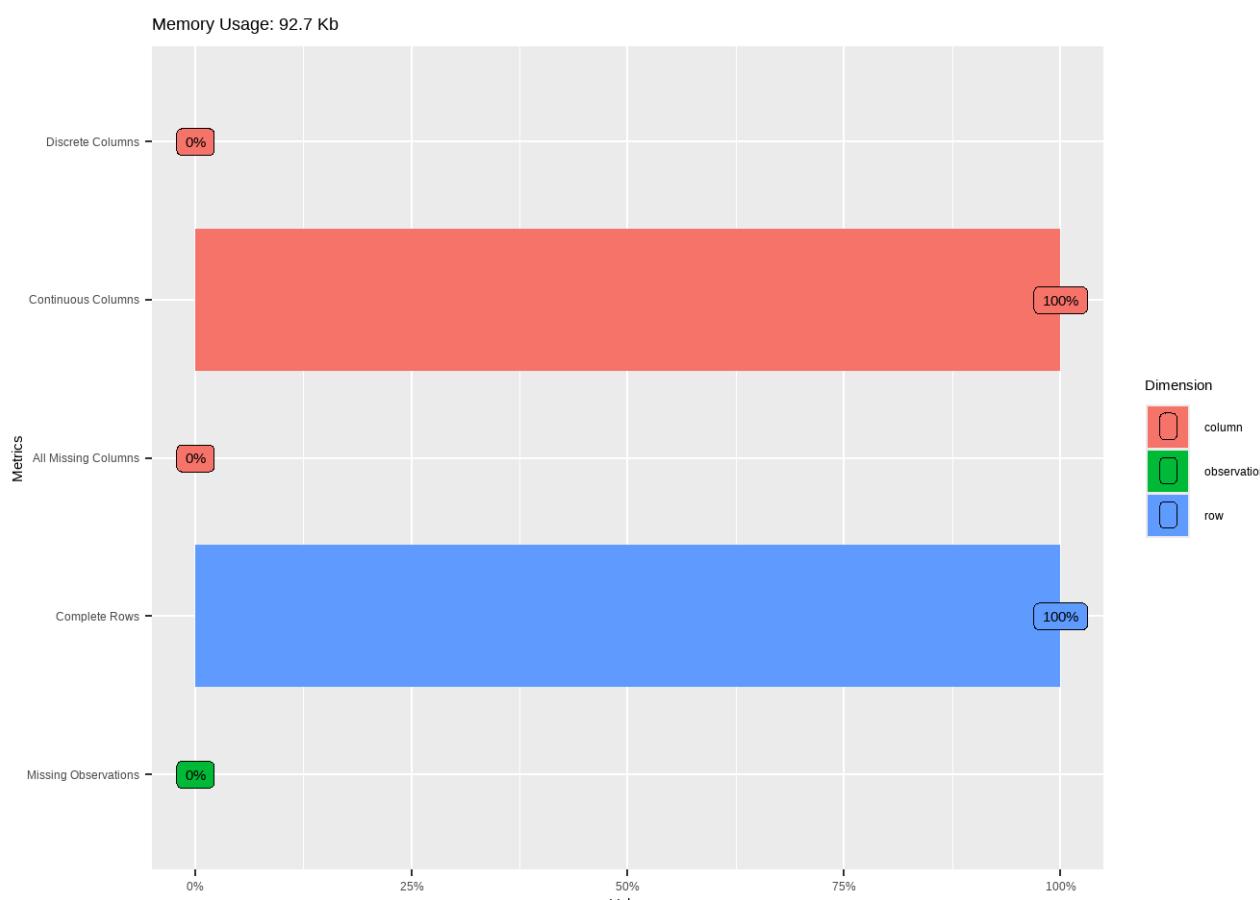
```
# A tibble: 6 × 25
  player_number   salary    age height weight games_played games_started
  <dbl>      <dbl> <dbl> <dbl>  <dbl>       <dbl>        <dbl>
1 1 947276 36 79 260 966 838
2 2 25000000 37 78 212 1346 1198
3 3 4088019 39 78 220 1274 954
4 4 5675000 36 77 195 1100 432
5 5 5250000 40 83 250 1392 1389
6 6 8500000 39 83 240 1462 1425
# i 18 more variables: minutes_per_game <dbl>, fg_made <dbl>,
#   fg_attempted <dbl>, fg_percent <dbl>, x3p_made <dbl>, x3p_attempted <dbl>,
#   x3p_percent <dbl>, ft_made <dbl>, ft_attempted <dbl>, ft_percent <dbl>,
```

```
# offensive_rebounds <dbl>, defensive_rebounds <dbl>, assists <dbl>,
# blocks <dbl>, steals <dbl>, personal_fouls <dbl>, turnovers <dbl>,
# points <dbl>
```

```
nba_data %>% tail()
```

```
# A tibble: 6 × 25
  player_number salary age height weight games_played games_started
  <dbl>    <dbl> <dbl> <dbl> <dbl>      <dbl>        <dbl>
1        440 5103120   20    77   195       80          48
2        441 1733040   19    78   202       70           6
3        442 1140240   21    83   200       24           4
4        443 1131960   20    81   220        5           0
5        444 3102240   20    77   200       68          66
6        445 525093    23    79   185       12           0
# i 18 more variables: minutes_per_game <dbl>, fg_made <dbl>,
# fg_attempted <dbl>, fg_percent <dbl>, x3p_made <dbl>, x3p_attempted <dbl>,
# x3p_percent <dbl>, ft_made <dbl>, ft_attempted <dbl>, ft_percent <dbl>,
# offensive_rebounds <dbl>, defensive_rebounds <dbl>, assists <dbl>,
# blocks <dbl>, steals <dbl>, personal_fouls <dbl>, turnovers <dbl>,
# points <dbl>
```

```
nba_data %>% plot_intro()
```

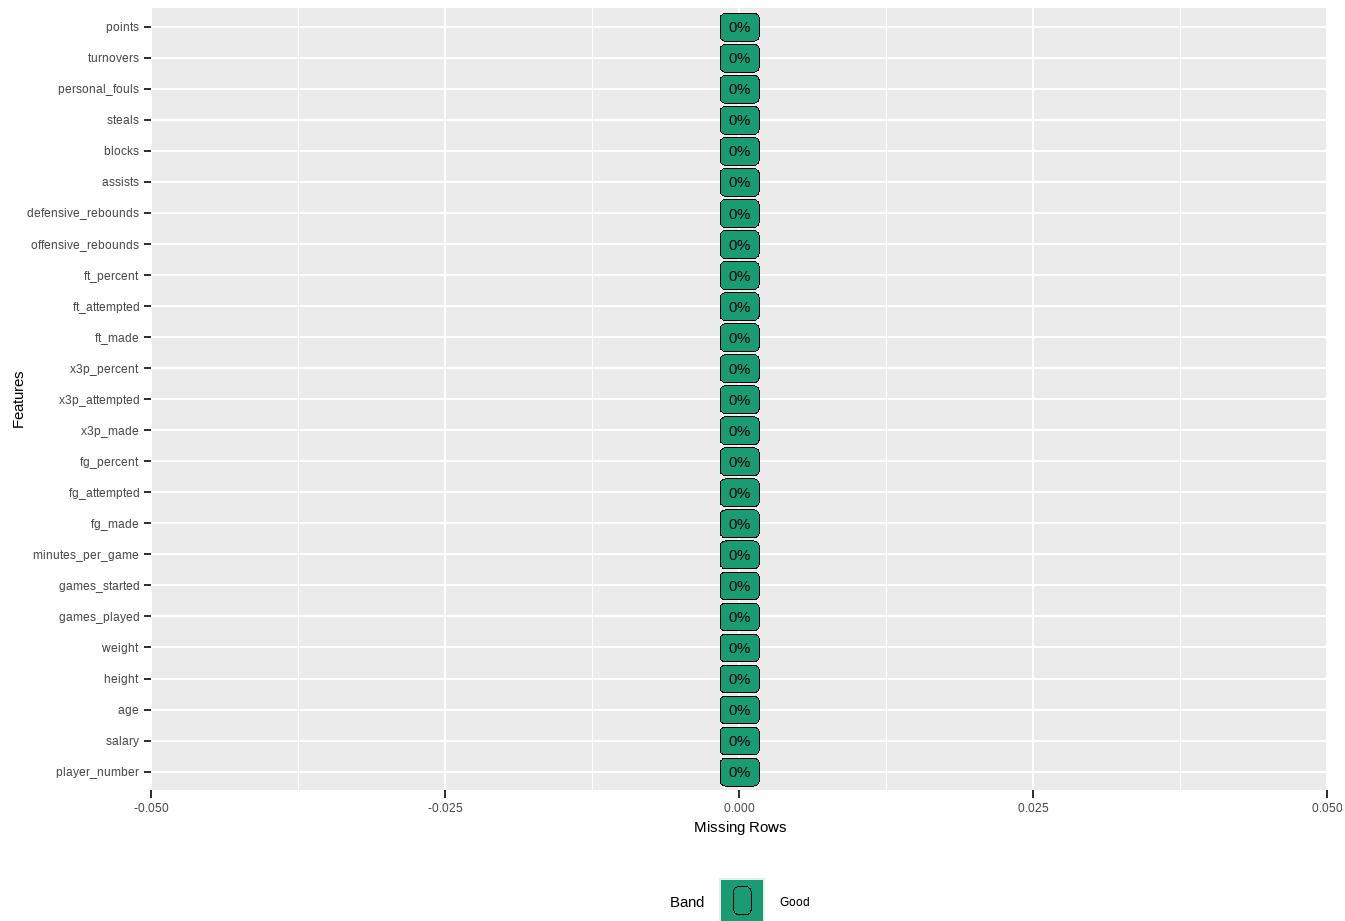


```
nba_data %>% glimpse()
```

Rows: 445
 Columns: 25

\$ player_number	<dbl> 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, ...
\$ salary	<dbl> 947276, 25000000, 4088019, 5675000, 5250000, 850000...
\$ age	<dbl> 36, 37, 39, 36, 40, 39, 38, 40, 36, 38, 37, 38, 38, ...
\$ height	<dbl> 79, 78, 78, 77, 83, 83, 78, 75, 80, 82, 84, 79, 74, ...
\$ weight	<dbl> 260, 212, 220, 195, 250, 240, 205, 200, 218, 250, 2...
\$ games_played	<dbl> 966, 1346, 1274, 1100, 1392, 1462, 923, 1304, 1012, ...
\$ games_started	<dbl> 838, 1198, 954, 432, 1389, 1425, 349, 939, 569, 354...
\$ minutes_per_game	<dbl> 32.4, 36.1, 32.4, 30.9, 34.0, 34.5, 26.3, 30.9, 27...
\$ fg_made	<dbl> 4.8, 8.7, 6.8, 5.4, 7.4, 7.2, 4.6, 4.6, 4.0, 2.4, 7...
\$ fg_attempted	<dbl> 11.6, 19.5, 15.5, 13.1, 14.6, 14.5, 10.1, 10.0, 8.6...
\$ fg_percent	<dbl> 0.415, 0.447, 0.439, 0.410, 0.506, 0.497, 0.450, 0...
\$ x3p_made	<dbl> 1.2, 1.4, 1.5, 1.8, 0.0, 0.1, 1.5, 0.1, 1.6, 0.0, 1...
\$ x3p_attempted	<dbl> 3.5, 4.1, 4.1, 5.0, 0.1, 0.4, 3.9, 0.7, 3.8, 0.0, 3...
\$ x3p_percent	<dbl> 0.340, 0.329, 0.373, 0.349, 0.179, 0.275, 0.369, 0...
\$ ft_made	<dbl> 2.7, 6.2, 3.7, 3.0, 4.2, 3.3, 3.4, 3.1, 1.3, 1.0, 5...
\$ ft_attempted	<dbl> 3.7, 7.4, 4.6, 3.5, 6.1, 4.2, 4.2, 3.8, 1.7, 1.6, 5...
\$ ft_percent	<dbl> 0.716, 0.837, 0.800, 0.861, 0.696, 0.789, 0.827, 0...
\$ offensive_rebounds	<dbl> 1.2, 1.1, 1.2, 0.4, 2.8, 2.2, 0.7, 1.1, 0.6, 1.8, 1...
\$ defensive_rebounds	<dbl> 3.4, 4.1, 3.4, 2.0, 8.1, 7.8, 3.0, 2.6, 3.7, 2.9, 6...
\$ assists	<dbl> 2.7, 4.7, 3.4, 3.6, 3.0, 3.7, 4.0, 6.5, 2.6, 0.4, 2...
\$ blocks	<dbl> 0.5, 0.5, 0.6, 0.2, 2.2, 1.4, 0.3, 0.2, 0.2, 0.6, 0...
\$ steals	<dbl> 1.8, 1.4, 1.1, 1.0, 0.7, 1.3, 1.4, 1.2, 0.6, 0.4, 0...
\$ personal_fouls	<dbl> 2.7, 2.5, 2.8, 1.6, 2.4, 2.4, 2.1, 2.2, 2.0, 2.1, 2...
\$ turnovers	<dbl> 1.8, 3.0, 1.9, 2.0, 2.4, 2.2, 2.1, 2.4, 1.5, 0.9, 1...
\$ points	<dbl> 13.5, 25.0, 18.8, 15.5, 19.0, 17.8, 14.0, 12.5, 10...

```
nba_data %>% plot_missing()
```



```
nba_data %>% str()
```

```
tibble [445 x 25] (S3:tbl_df/tbl/data.frame)
$ player_number : num [1:445] 1 2 3 4 5 6 7 8 9 10 ...
$ salary        : num [1:445] 947276 25000000 4088019 5675000 5250000 ...
$ age          : num [1:445] 36 37 39 36 40 39 38 40 36 38 ...
$ height        : num [1:445] 79 78 78 77 83 83 78 75 80 82 ...
$ weight        : num [1:445] 260 212 220 195 250 240 205 200 218 250 ...
$ games_played : num [1:445] 966 1346 1274 1100 1392 ...
$ games_started: num [1:445] 838 1198 954 432 1389 ...
$ minutes_per_game: num [1:445] 32.4 36.1 32.4 30.9 34 34.5 26.3 30.9 27.3 15.8 ...
$ fg_made       : num [1:445] 4.8 8.7 6.8 5.4 7.4 7.2 4.6 4.6 4 2.4 ...
$ fg_attempted  : num [1:445] 11.6 19.5 15.5 13.1 14.6 14.5 10.1 10 8.6 4.9 ...
$ fg_percent    : num [1:445] 0.415 0.447 0.439 0.41 0.506 0.497 0.45 0.461 0.459 0.486 ...
$ x3p_made     : num [1:445] 1.2 1.4 1.5 1.8 0 0.1 1.5 0.1 1.6 0 ...
$ x3p_attempted: num [1:445] 3.5 4.1 4.1 5 0.1 0.4 3.9 0.7 3.8 0 ...
$ x3p_percent   : num [1:445] 0.34 0.329 0.373 0.349 0.179 0.275 0.369 0.217 0.407 0 ...
$ ft_made       : num [1:445] 2.7 6.2 3.7 3 4.2 3.3 3.4 3.1 1.3 1 ...
$ ft_attempted  : num [1:445] 3.7 7.4 4.6 3.5 6.1 4.2 4.2 3.8 1.7 1.6 ...
$ ft_percent    : num [1:445] 0.716 0.837 0.8 0.861 0.696 0.789 0.827 0.807 0.769 0.64 ...
$ offensive_rebounds: num [1:445] 1.2 1.1 1.2 0.4 2.8 2.2 0.7 1.1 0.6 1.8 ...
$ defensive_rebounds: num [1:445] 3.4 4.1 3.4 2 8.1 7.8 3 2.6 3.7 2.9 ...
$ assists       : num [1:445] 2.7 4.7 3.4 3.6 3 3.7 4 6.5 2.6 0.4 ...
```

```
$ blocks      : num [1:445] 0.5 0.5 0.6 0.2 2.2 1.4 0.3 0.2 0.2 0.6 ...
$ steals      : num [1:445] 1.8 1.4 1.1 1 0.7 1.3 1.4 1.2 0.6 0.4 ...
$ personal_fouls : num [1:445] 2.7 2.5 2.8 1.6 2.4 2.4 2.1 2.2 2 2.1 ...
$ turnovers   : num [1:445] 1.8 3 1.9 2 2.4 2.2 2.1 2.4 1.5 0.9 ...
$ points     : num [1:445] 13.5 25 18.8 15.5 19 17.8 14 12.5 10.8 5.8 ...
```

```
nba_data %>% ncol()
```

[1] 25

```
nba_data %>% nrow()
```

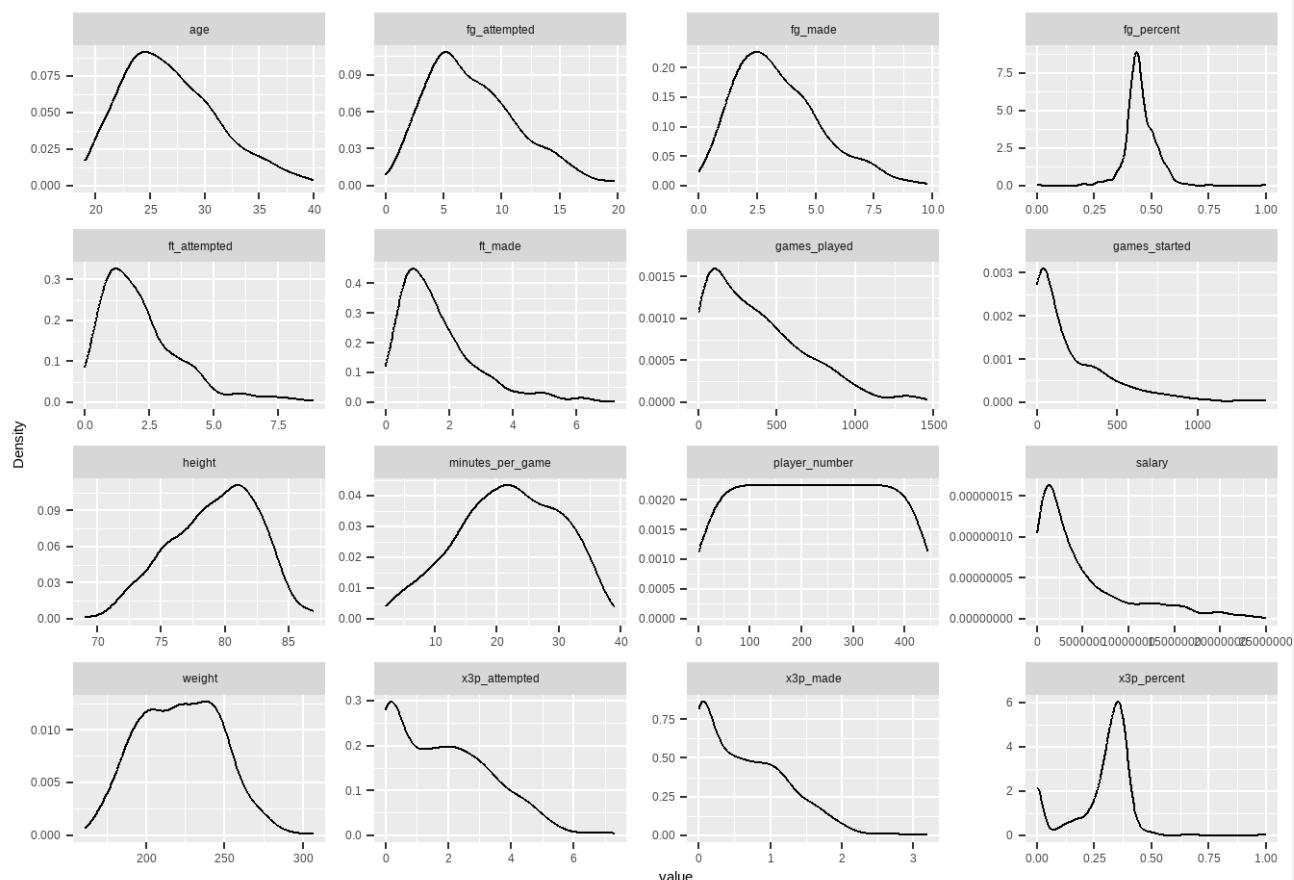
[1] 445

Comments on Dataset Exploration:

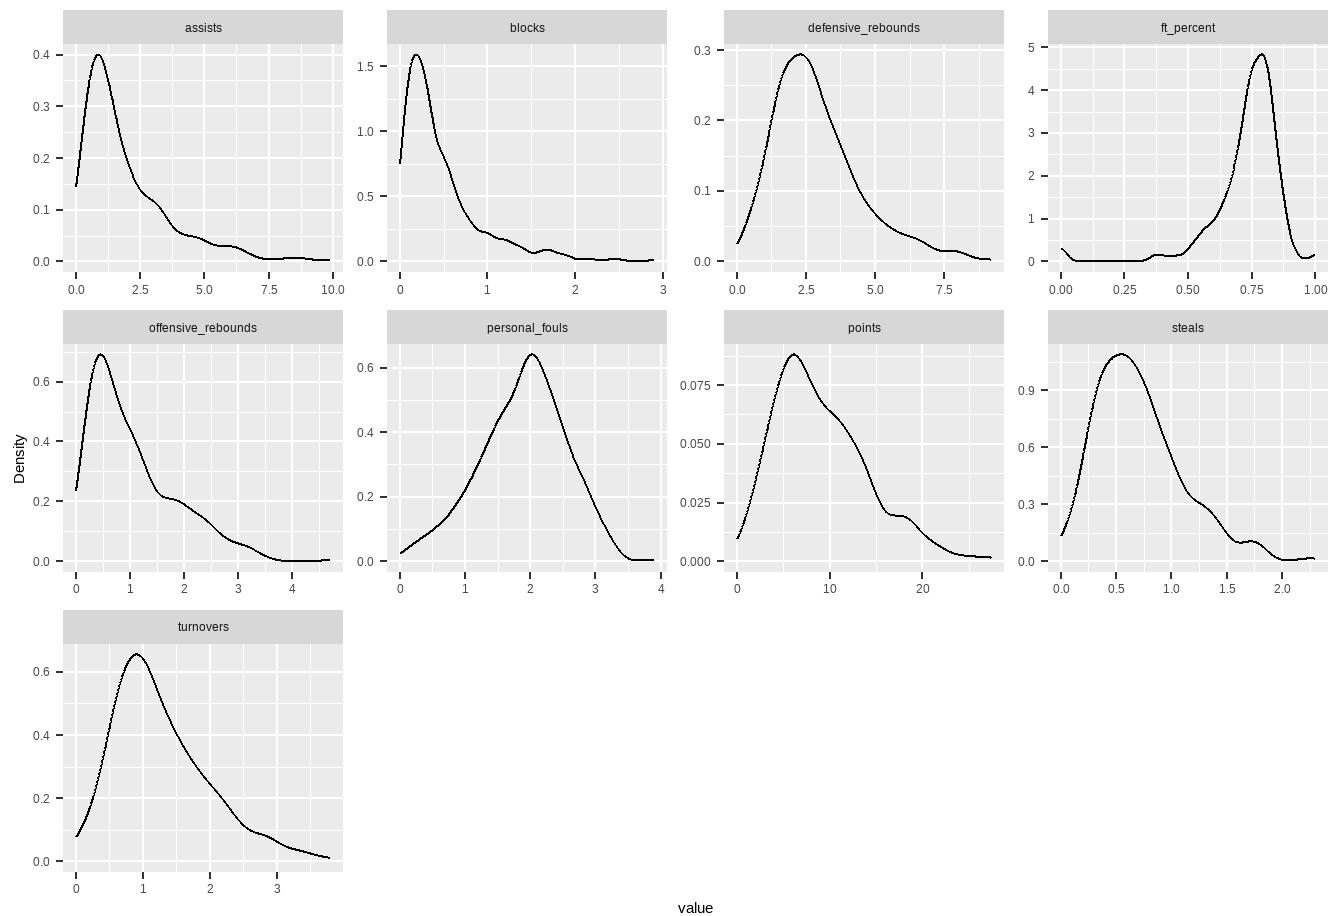
This dataset is made of only numeric variables, contains no categorical variables. This dataset also contains no missing values.

Variable EDA

```
nba_data %>% plot_density()
```



Page 1

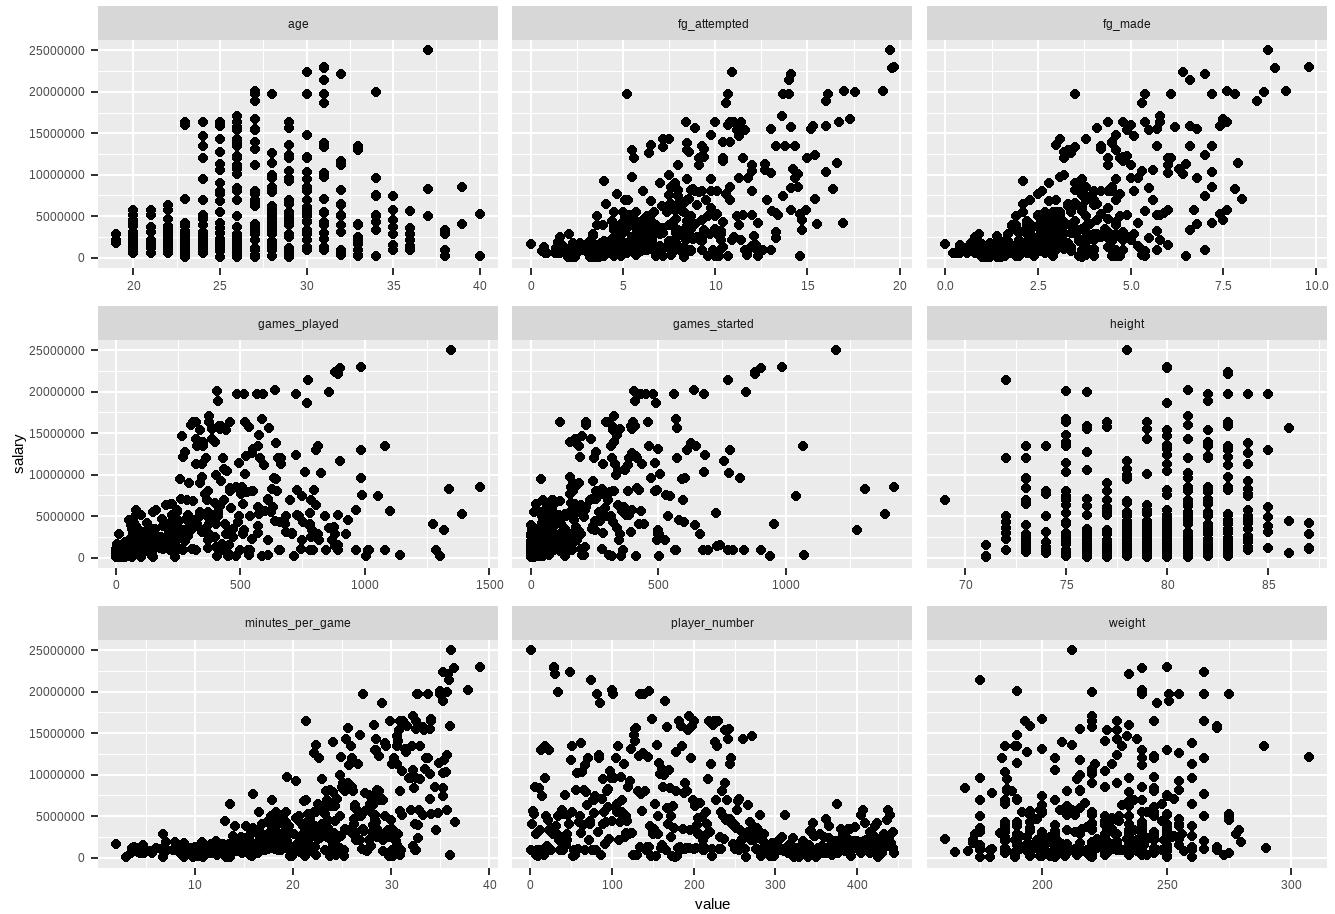


Page 2

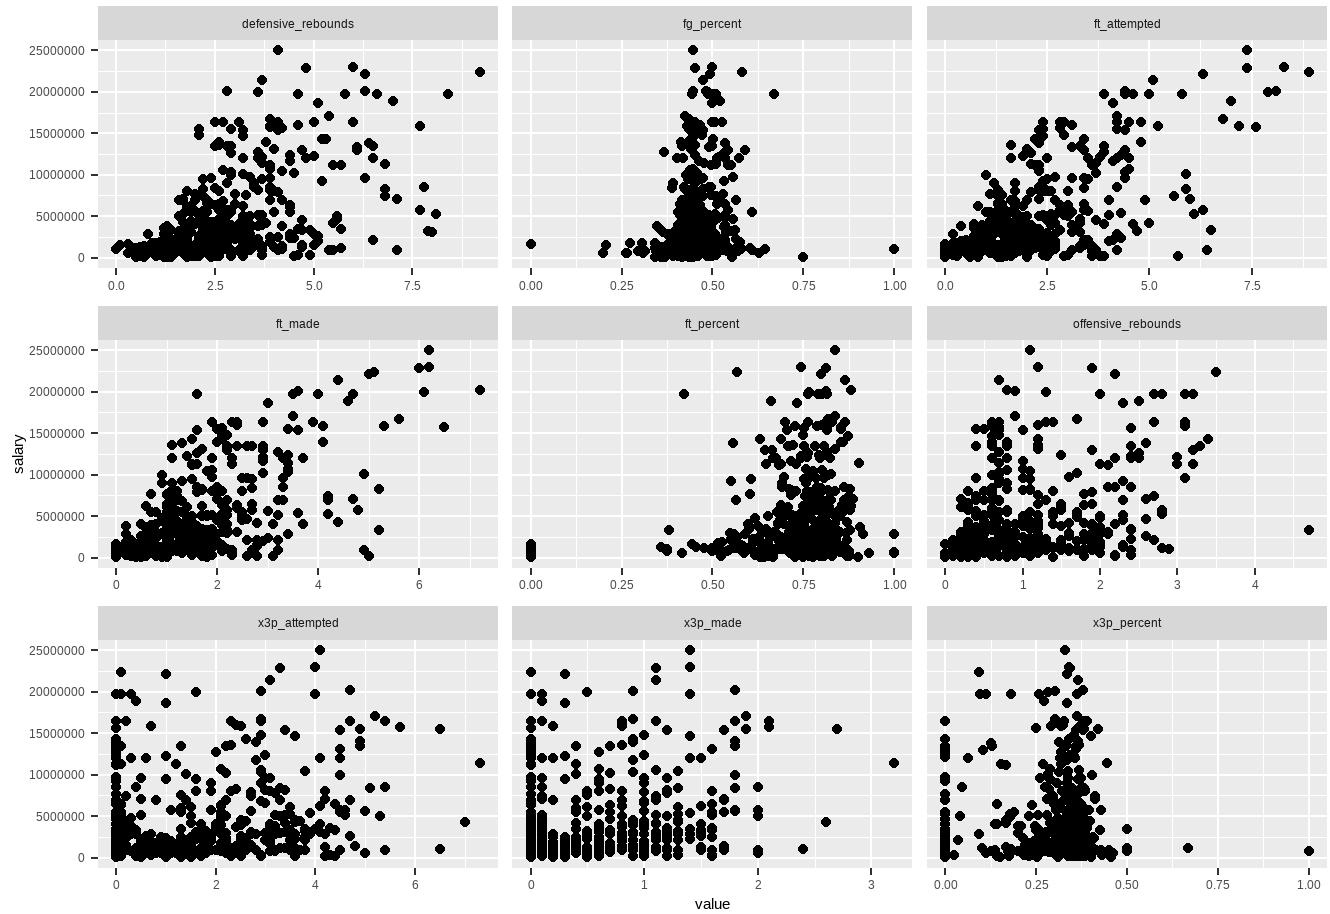
```
diagnose_outlier(nba_data) ## Although there are outlier counts in almost all variables it would
not be ideal to remove people that are outliers in categories because that could
determine salary, ex if a person is an outlier because he has more blocks but block
influence salary that player would earn more salary.
```

```
# A tibble: 25 × 6
  variables    outliers_cnt outliers_ratio outliers_mean with_mean without_mean
  <chr>        <int>      <dbl>       <dbl>      <dbl>      <dbl>
1 player_number...     0          0        NaN     2.23e2      223
2 salary            35        7.87    18164924.   4.84e6  3705946.
3 age                2        0.449       40     2.69e1      26.8
4 height              1        0.225       69     7.92e1      79.2
5 weight              1        0.225      307     2.22e2      221.
6 games_played         8        1.80     1340.     3.70e2      352.
7 games_start...       20        4.49    1015.     2.12e2      174.
8 minutes_per...       0          0        NaN     2.19e1      21.9
9 fg_made              6        1.35       8.93     3.41e0      3.33
10 fg_attempted         4        0.899      19.5     7.52e0      7.41
# i 15 more rows
```

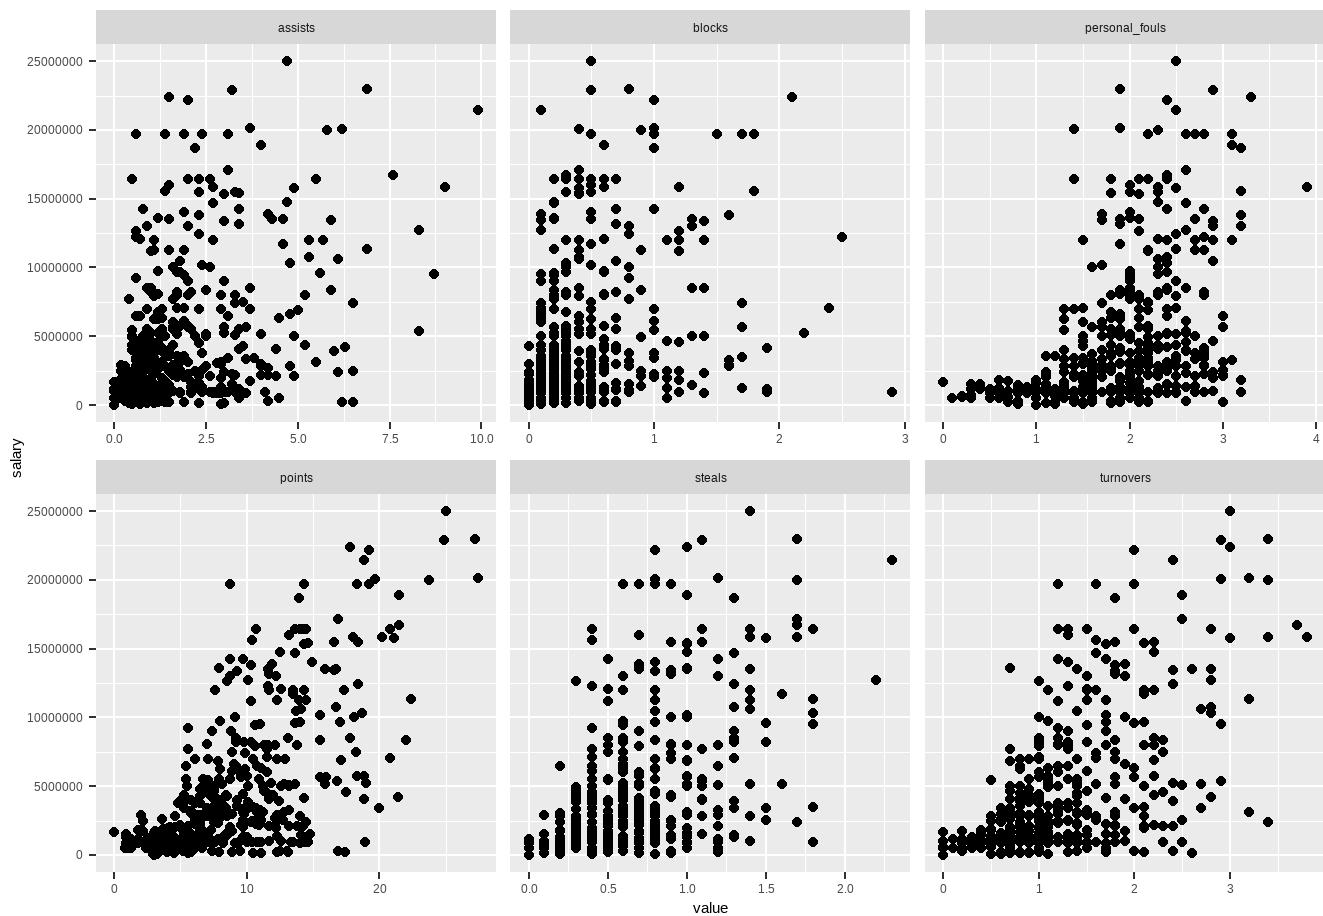
```
nba_data %>% plot_scatterplot(by="salary")
```



Page 1



Page 2



Page 3

```
skewness_values <- sapply(nba_data, skewness, na.rm = TRUE)
skew_table <- data.frame(
  Variable = names(skewness_values),
  Skewness = round(skewness_values, 3)
)
print(skew_table) ## Prints skew values of all numeric variables
```

	Variable	Skewness
player_number	player_number	0.000
salary	salary	1.569
age	age	0.601
height	height	-0.301
weight	weight	0.122
games_played	games_played	0.981
games_started	games_started	1.810
minutes_per_game	minutes_per_game	-0.215
fg_made	fg_made	0.738
fg_attempted	fg_attempted	0.648
fg_percent	fg_percent	0.556
x3p_made	x3p_made	0.845
x3p_attempted	x3p_attempted	0.642
x3p_percent	x3p_percent	-0.697
ft_made	ft_made	1.502

ft_attempted	ft_attempted	1.398
ft_percent	ft_percent	-2.777
offensive_rebounds	offensive_rebounds	1.121
defensive_rebounds	defensive_rebounds	1.063
assists	assists	1.668
blocks	blocks	1.918
steals	steals	0.890
personal_fouls	personal_fouls	-0.308
turnovers	turnovers	0.881
points	points	0.829

Simple Random Forest To Find Variables with low importance

```
rf_model <- randomForest(salary ~ ., data = nba_data, importance = TRUE)
imp <- importance(rf_model)

imp_sorted <- imp[order(imp[, "%IncMSE"]), , drop = FALSE]

# if you want both columns
imp_table <- data.frame(
  Variable      = rownames(imp_sorted),
  `%IncMSE`     = imp_sorted[, "%IncMSE"],
  IncNodePurity = imp_sorted[, "IncNodePurity"],
  row.names = NULL,
  check.names = FALSE
)
print(imp_table)
```

	Variable	%IncMSE	IncNodePurity
1	x3p_attempted	1.368133	104998860067654
2	weight	2.359012	153857781823841
3	height	2.589529	100537444272033
4	x3p_percent	2.790541	139173589374633
5	ft_percent	3.496960	185432723489313
6	assists	3.714161	137649255719861
7	turnovers	4.388886	230125973562826
8	x3p_made	4.403810	78264698949576
9	personal_fouls	5.449440	163991577429311
10	steals	5.807849	140729659538850
11	ft_attempted	6.799977	836759317636665
12	fg_percent	7.406890	266522088603520
13	blocks	7.486030	191505252733787
14	defensive_rebounds	7.887766	355276626549753
15	ft_made	8.034190	766716840070486
16	offensive_rebounds	9.128378	298226729291832
17	fg_made	9.202119	770673532152055
18	fg_attempted	9.391082	433064503004834
19	points	10.291604	1120464340915471
20	minutes_per_game	14.179956	1426834891603088

```

21      age 16.430643 526371897102351
22      player_number 17.654092 734368361264811
23      games_played 18.157244 742989010145924
24      games_started 18.636092 2057591595861610

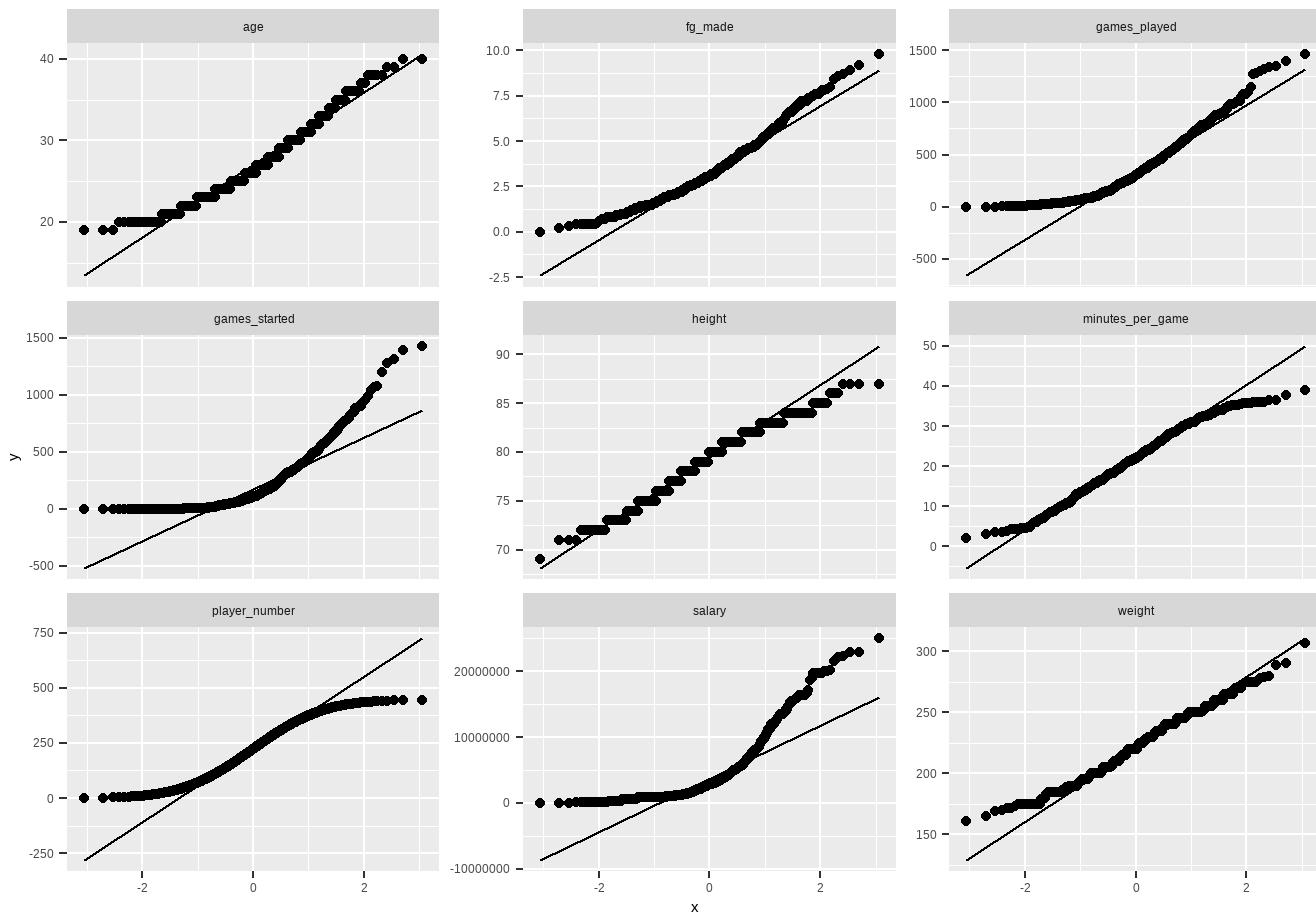
```

Comments on Simple Random Forest:

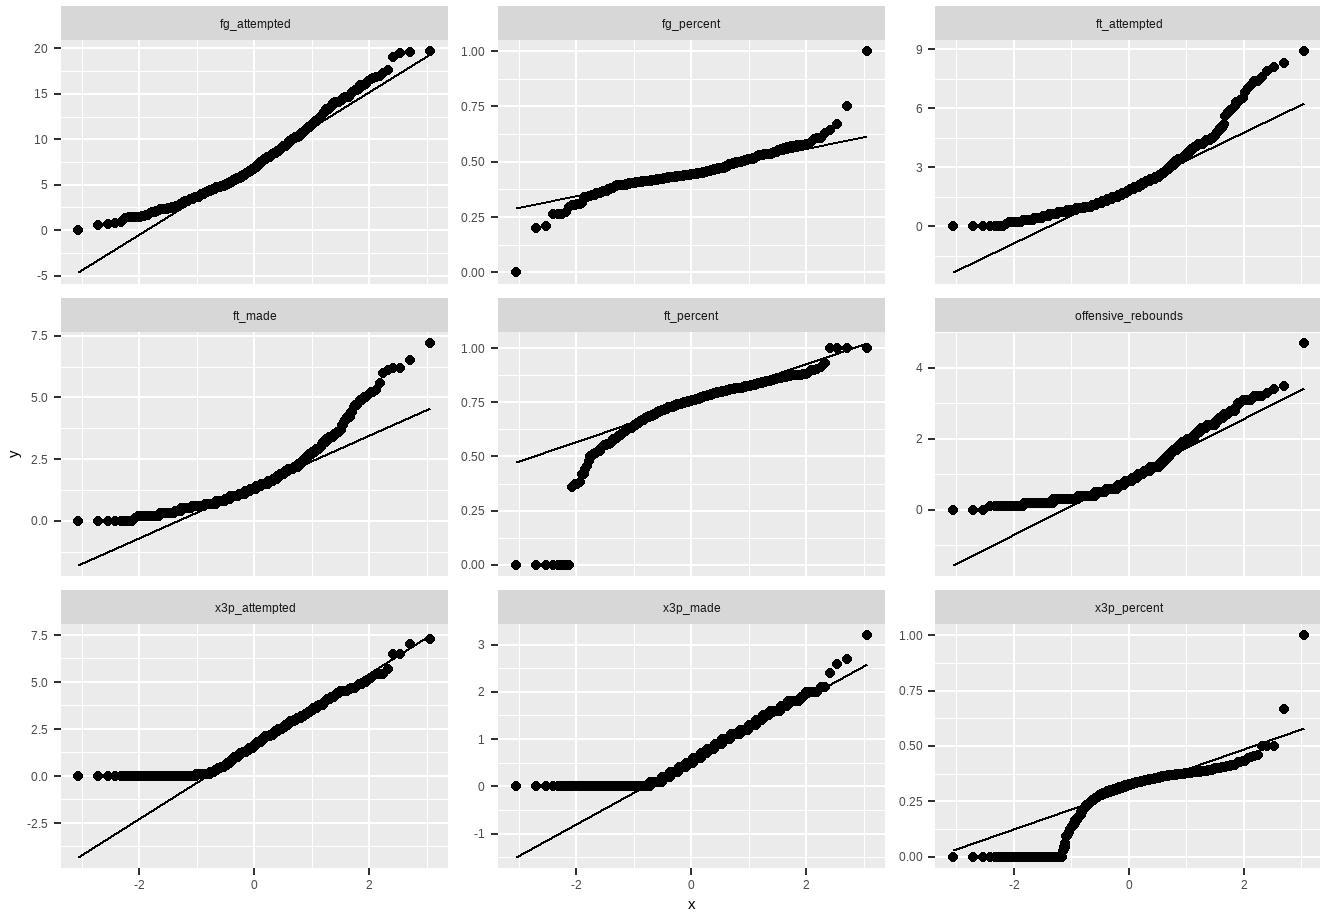
The variable importance feature of the random forest can help initially identify variables that are not important when compared to salary, the dependent variable. This helps get insight into potential variables to drop due to their lack of importance. Any variable that has a value of <5 for %IncMSE can be safely dropped.

Q-Q Plots

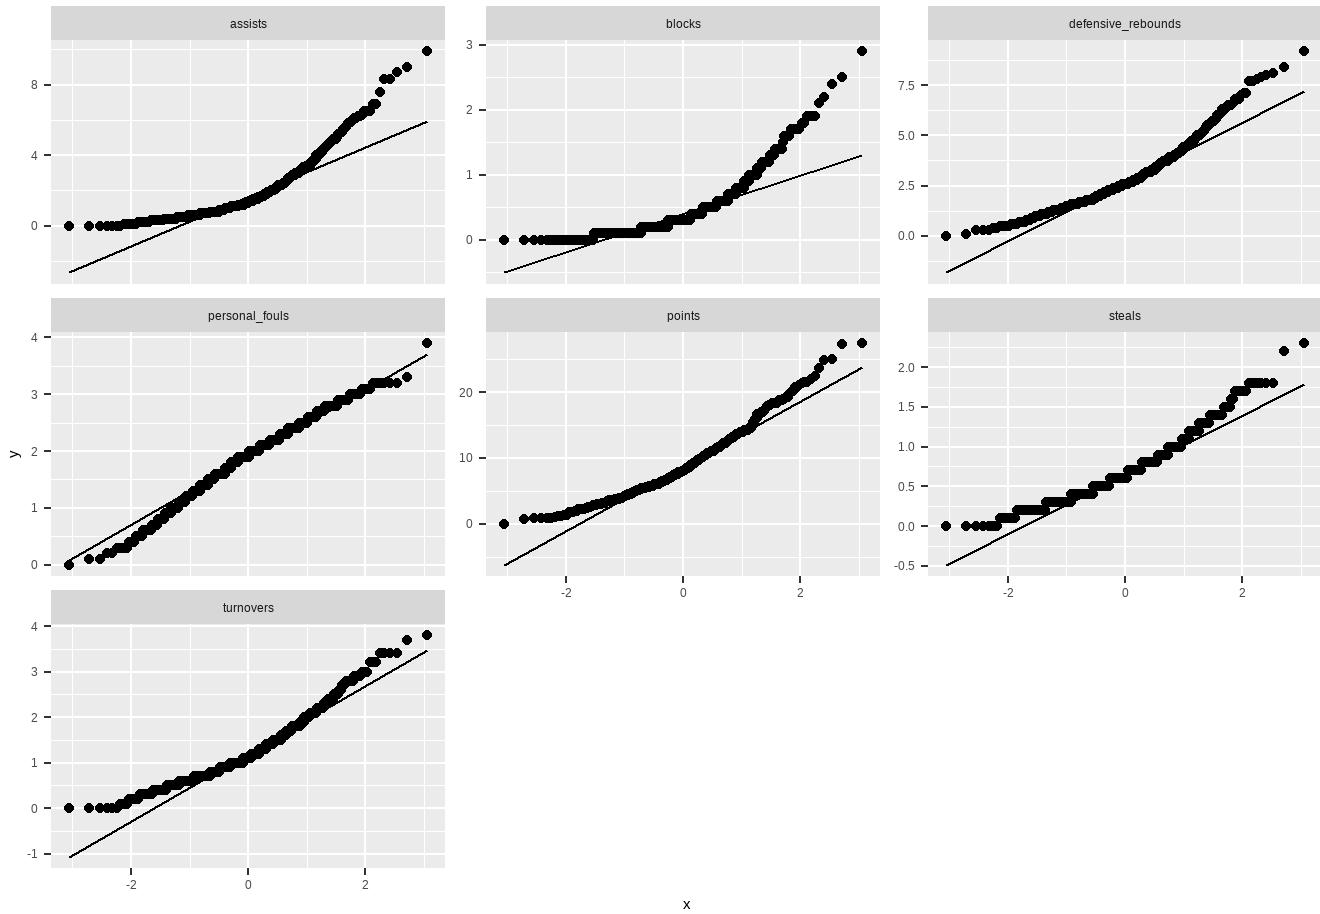
```
nba_data %>% plot_qq()
```



Page 1



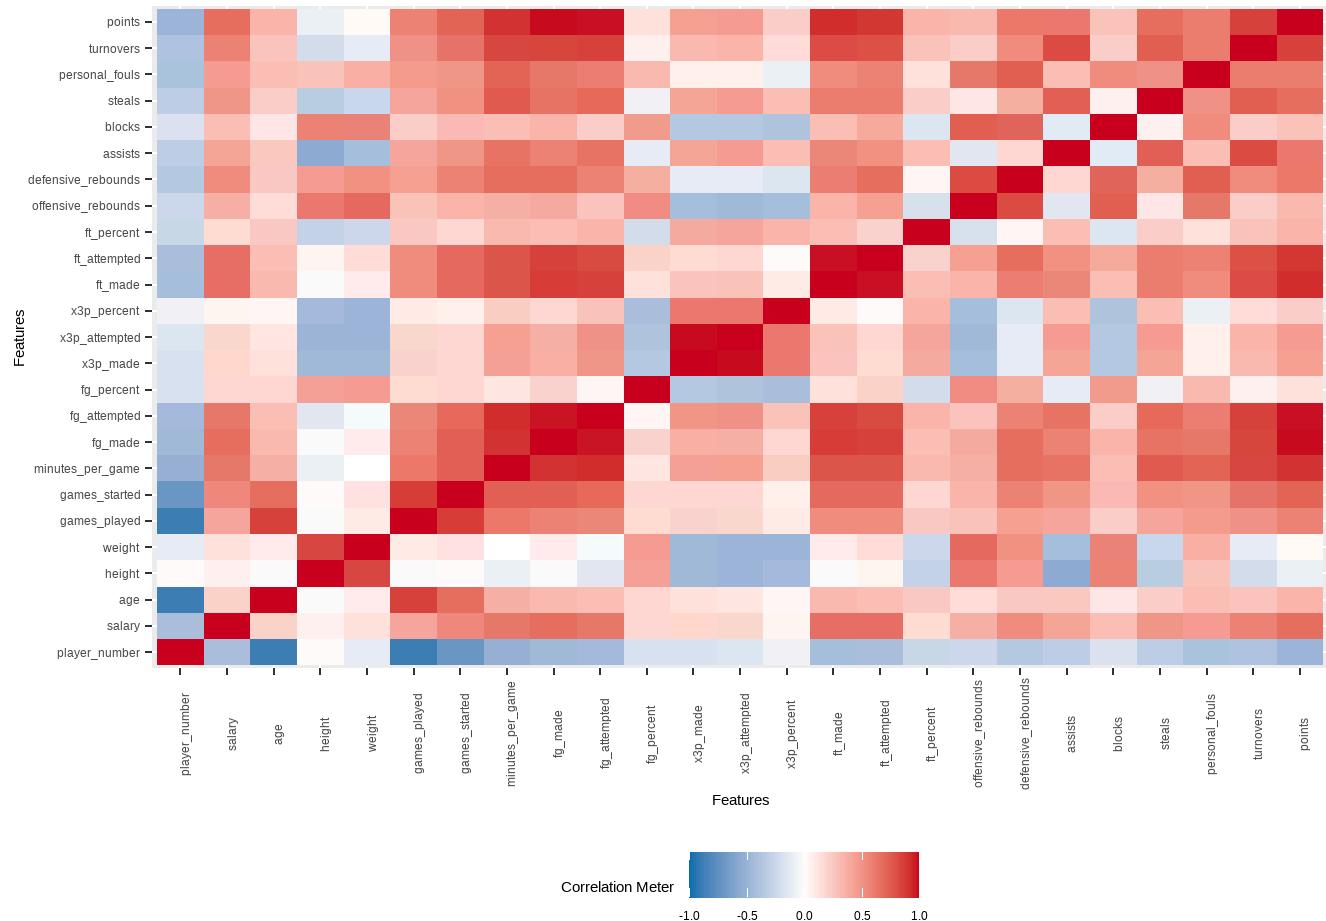
Page 2



Page 3

Correlation Matrix

```
nba_data %>% plot_correlation()
```



Data Preparation

Drop Player Number as a predictor

```
nba_data <- nba_data %>%
  select(-player_number)
```

Partition

```
myindex <- createDataPartition(nba_data$salary, p=0.7, list = FALSE)
trainSet <- nba_data[myindex,]
testSet <- nba_data[-myindex,]

cat("Mean of Dependent Variable \n")
```

Mean of Dependent Variable

```
mean(nba_data$salary)
```

```
[1] 4843169
```

```
cat("Mean of trainset dependent variable \n")
```

Mean of trainset dependent variable

```
mean(trainSet$salary)
```

```
[1] 4814817
```

```
cat("Mean of testset dependent variable \n")
```

Mean of testset dependent variable

```
mean(testSet$salary)
```

```
[1] 4910398
```

Comments on Data Partitioning:

The modeling dataset was partitioned into a 70/30 split, 70% of the data will be used to train the model and 30% of the dataset will be used to test the model. The means of salary for the train and test dataset were both close which means both have similar distributions of the dependent variable.

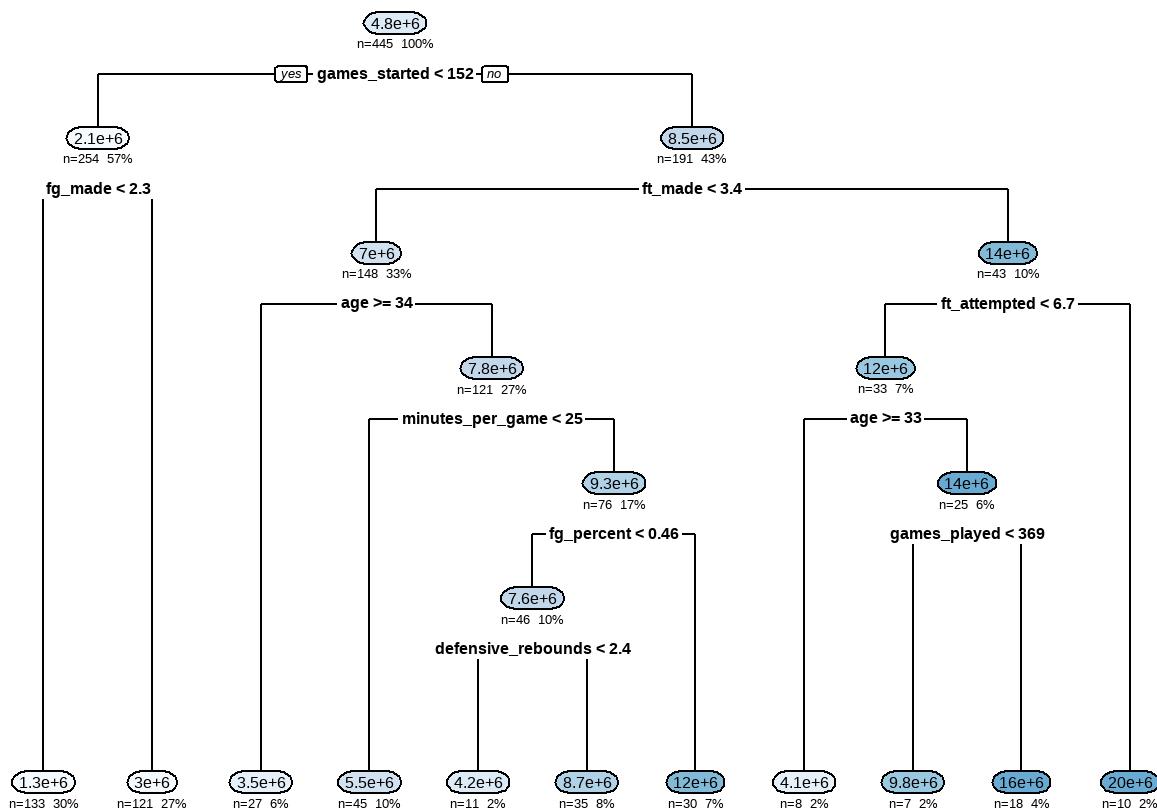
Modeling (Rpart)

Cross Validation

```
myCtrl <- trainControl(method = "cv", number = 10)
```

Default Tree (Rpart)

```
set.seed(1)
default_tree <- rpart(salary ~., data = nba_data, method = "anova")
rpart.plot(default_tree, type = 2, extra = 101, under = TRUE, fallen.leaves = TRUE)
```



Answer to Question 45 Part A:

A) What are the predictor variable and split value for the first split of the default regression tree?

The Predictor Variable for the first split is games started. The split values are for people that have less than 152 games started their salary is 2,100,000 whereas people with more than 152 games started their salary is 8,500,000.

Full Tree (Rpart)

```
set.seed(1)
full_tree <- rpart(salary ~., data = nba_data, control = rpart.control(cp = 0, minsplit = 2,
minbucket = 1, maxdepth = 30, xval = 10))

#xval = 10 is cross fold validation of 10, this is how to do it using rpart, used rpart to get
correct values in the cp table

head(full_tree$cptable)
```

	CP	nsplit	rel error	xerror	xstd
1	0.36913609	0	1.0000000	1.0046536	0.09285005
2	0.11584050	1	0.6308639	0.6867124	0.06355926
3	0.04723599	2	0.5150234	0.6693572	0.06324465
4	0.03493893	4	0.4205514	0.6011018	0.06073049

```
5 0.03326144      5 0.3856125 0.5903899 0.06014571
6 0.02580501      6 0.3523511 0.5424423 0.05800483
```

```
## Where to Find Answers to Part B
```

```
min_error <- full_tree$cptable[which.min(full_tree$cptable[, "xerror"]), ]
min_error
```

CP	nsplit	rel error	xerror	xstd
0.01786872	7.0000000	0.32654605	0.53163737	0.06063138

```
## Find number of Leaf nodes in the minimum error tree
num_leaves <- min_error["nspit"]
cat("Number of Leaf Nodes in the Minimum Tree: ")
```

Number of Leaf Nodes in the Minimum Tree:

```
print(num_leaves)
```

```
nspit
7
```

Answer to Question 45 Part B:

B) Build a full-grown tree. Which cp value is associated with the lowest cross-validation error? How many leaf nodes are in the minimum-error tree?

The CP value associated with the lowest cross-validation error is 0.01786872. There are seven leaf nodes in the minimum-error tree.

Best Pruned Tree

```
cpt <- full_tree$cptable

imin <- which.min(cpt[, "xerror"]) ## Grabs the row that the minimum error is on and defines it
with a variable
xerr_min <- cpt[imin, "xerror"] ## Gets the xerror
xstd_min <- cpt[imin, "xstd"] ## Gets the xstd
threshold <- xerr_min + xstd_min ## Creates the xerror threshold for best pruned tree.
threshold ## lowest amount of splits while staying below threshold xerror value
```

```
[1] 0.5922688
```

```
within_threshold <- which(cpt[, "xerror"] <= threshold)

best_pruned_index <- within_threshold[which.min(cpt[within_threshold, "nspit"])]
```

```
best_pruned_cp <- cpt[best_pruned_index, "CP"]
```

```
best_pruned_xerror <- cpt[best_pruned_index, "xerror"]
best_pruned_nsplits <- cpt[best_pruned_index, "nsplit"]

cat("Best-Pruned Tree (1-SE Rule)\n")
```

Best-Pruned Tree (1-SE Rule)

```
cat("CP value:", best_pruned_cp, "\n")
```

CP value: 0.03326144

```
cat("Cross-validation error:", best_pruned_xerror, "\n")
```

Cross-validation error: 0.5903899

```
cat("Number of splits:", best_pruned_nsplits, "\n")
```

Number of splits: 5

Answer to Question 45 Part C:

C) Is there a simpler tree with a cross-validation error that is within one standard error of the minimum error? If there is, then which cp value is associated with the best-pruned tree?

Yes, there is a simpler tree with a cross-validation error that is within one standard error of the minimum error. This tree has a cp value of 0.03326144.

Best Tree

```
best_pruned_tree <- prune(full_tree, cp=0.03326144)
```

Variable Importance

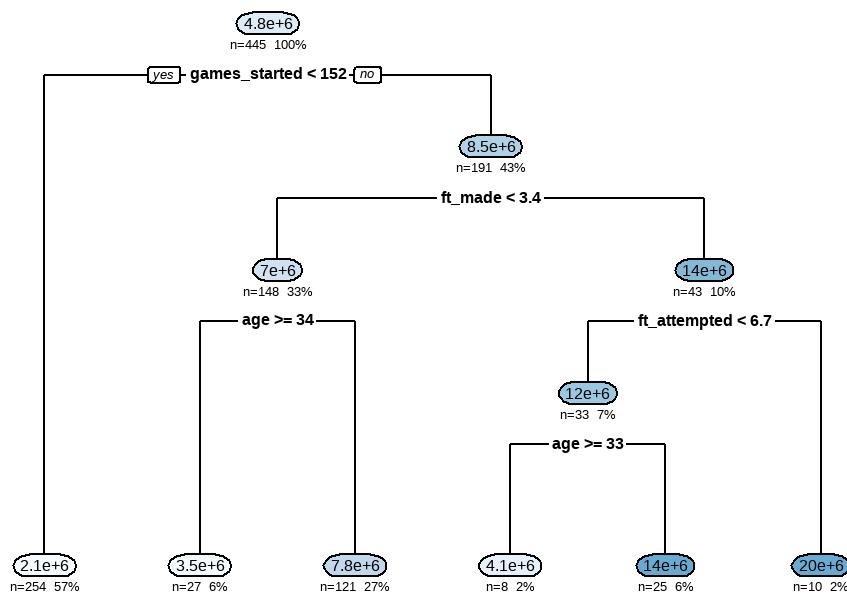
```
caret::varImp(best_pruned_tree)
```

	Overall
age	0.78148610
defensive_rebounds	0.07855415
fg_attempted	0.23861331
fg_made	0.72078128
ft_attempted	0.48212458
ft_made	0.50742141
games_played	0.31387854
games_started	0.51839552
minutes_per_game	0.69766452
offensive_rebounds	0.09162487
personal_fouls	0.08197401

points	0.73324452
x3p_percent	0.16038949
height	0.00000000
weight	0.00000000
fg_percent	0.00000000
x3p_made	0.00000000
x3p_attempted	0.00000000
ft_percent	0.00000000
assists	0.00000000
blocks	0.00000000
steals	0.00000000
turnovers	0.00000000

Show Best Pruned Tree

```
rpart.plot(best_pruned_tree, type = 2, extra = 101, under = TRUE, fallen.leaves = TRUE)
```



```
cat("Number of Leaf Nodes in the Default Tree: ")
```

Number of Leaf Nodes in the Default Tree:

```
sum(best_pruned_tree$frame$var == "<leaf>")
```

[1] 6

Answer to Question 45 Part D:

- D. Prune the full tree to the best-pruned tree or the minimum-error tree if the answer to part (c) is "No." Display the tree. What are the rules that can be derived from the pruned tree?

From the best pruned tree, it can be derived that games started, ft made, age, and ft attempted are important predictor variables.

- Path 1: Players who started fewer than 152 games have an average predicted salary of about \$2.1 million.
- Path 2: Players who started 152 or more games, made fewer than 3.4 free throws, and are aged 34 or older have an average salary of about \$3.5 million.
- Path 3: Players who started 152 or more games, made fewer than 3.4 free throws, and are younger than 34 have an average salary of about \$7.8 million.
- Path 4: Players who started 152 or more games, made at least 3.4 free throws, attempted fewer than 6.7 free throws, and are aged 33 or older have an average salary of about \$4.1 million.
- Path 5: Players who started 152 or more games, made at least 3.4 free throws, attempted fewer than 6.7 free throws, and are younger than 33 have an average salary of about \$12 million.
- Path 6: Players who started 152 or more games, made at least 3.4 free throws, and attempted 6.7 or more free throws have the highest predicted salary of about \$20 million.

Evaluation (Rpart)

Evaluation Metrics

```

predict_bpt <- predict(best_pruned_tree, testSet)
bpt_metrics <- round(forecast::accuracy(predict_bpt, testSet$salary),2)

## For inline code
ME_value   <- bpt_metrics["Test set", "ME"]
RMSE_value <- bpt_metrics["Test set", "RMSE"]
MAE_value  <- bpt_metrics["Test set", "MAE"]
MPE_value  <- bpt_metrics["Test set", "MPE"]
MAPE_value <- bpt_metrics["Test set", "MAPE"]

rownames(bpt_metrics) <- "Rpart Regression Tree"
print(bpt_metrics)

```

	ME	RMSE	MAE	MPE	MAPE
Rpart Regression Tree	-150347.9	3210484	2345165	-231.71	256.97

Answer to Question 45 Part E:

E. What are the ME, RMSE, MAE, MPE, and MAPE of the pruned tree on the validation data?

- **ME Value:** 'r ME_value'.
- **RMSE Value:** 'r RMSE_value'.
- **MAE Value:** 'r MAE_value'.
- **MPE Value:** 'r MPE_value' %.
- **MAPE Value:** 'r MAPE_value' %.

Measures of Accuracy:

- RMSE (Root Mean Square Error = 894.68) and MAE (Mean Absolute Error = 713.39) indicate that on average, the pruned tree's salary predictions deviate by roughly 700–900 from the true values.
- MAPE (Mean Absolute Percentage Error = 34.74%) means predictions are off by about 35% on average relative to the actual salary level.
- Overall, the model shows moderate prediction accuracy, capturing general salary trends but still leaving room for improvement in precision.

Measures of Bias:

- ME (Mean Error = -1.35) is extremely close to zero, showing the model is essentially unbiased — there is no meaningful tendency to systematically over- or under-predict salary.
- MPE (Mean Percentage Error = -14.75%) is slightly negative, indicating a mild tendency to under-predict salaries on average, but the bias is not large relative to the scale of the values.

Score (Rpart)

```
rpart_score <- predict(best_pruned_tree, nba_score)
rpart_pip <- cbind(nba_score, rpart_score)
rpart_pip
```

	player	age	height	weight	games_played	games_started	minutes_per_game	fg_made
1	Player1	36	81	260	984	820	31.8	6.8
2	Player2	32	76	185	714	504	30.2	5.4
3	Player3	27	85	235	636	472	32.0	5.2
	fg_attempted	fg_percent	x3p_made	x3p_attempted	x3p_percent	ft_made		
1	14.3	0.473	0.1	0.5	0.260	3.3		
2	11.0	0.437	1.6	4.2	0.384	5.0		
3	11.8	0.438	1.0	3.0	0.339	2.6		
	ft_attempted	ft_percent	offensive_rebounds	defensive_rebounds	assists	blocks		
1	4.4	0.765		3.1	6.3	1.8	0.3	
2	5.7	0.870		0.6	2.6	2.1	0.1	
3	3.4	0.786		1.0	3.9	1.7	0.6	
	steals	personal_fouls	turnovers	points	rpart_score			
1	0.8		2.4	2.1	18.0	3454649		
2	0.9		1.9	1.7	16.5	13977446		
3	0.8		2.4	1.7	15.0	7843275		

Answer to Question 45 Part F:

F) Score the three NBA players Merrick is trying to sign as ACE Sports Management clients in the NBA_Score worksheet using the pruned tree. What is the average predicted salary of the three players?

- Player 1: The Average salary for Player 1 was 3,454,649.
- Player 2: The Average salary for Player 2 was 13,977,446.
- Player 3: The Average salary for Player 3 is 7,843,275.

Modeling (Caret)

Best Tree (Caret)

```
caret_bp <- train(salary~., data = trainSet,
                  method = "rpart",
                  trControl = myCtrl,
                  tuneLength = 25,
                  metric = "RMSE",
                  control = rpart::rpart.control(minsplit = 2, minbucket = 1, cp = 0))

caret_bp$resample
```

	RMSE	Rsquared	MAE	Resample
1	4120411	0.266250283	3023060	Fold01
2	2995713	0.510020593	2212342	Fold07
3	4109285	0.378319536	3241325	Fold08
4	4078978	0.474508720	2778811	Fold05
5	5373464	0.004145221	3786076	Fold04
6	3970711	0.391088715	3059824	Fold06
7	3357730	0.425822028	2351667	Fold03
8	4799444	0.182573225	3247050	Fold02
9	4628543	0.360334374	3420274	Fold10
10	4719672	0.452011479	2873042	Fold09

CP Values

```
caret_bp$results
```

	cp	RMSE	Rsquared	MAE	RMSESD	RquaredSD	MAESD
1	0.00000000	4883057	0.3040928	3096999	695785.0	0.1439206	357985.9
2	0.01653605	4639461	0.3209965	3000470	611877.9	0.1522815	385906.4
3	0.03307209	4270504	0.3443940	2964679	452755.4	0.1333859	361348.6
4	0.04960814	4223961	0.3526245	2930154	507358.0	0.1340030	408134.3
5	0.06614419	4252855	0.3421655	2941652	522804.0	0.1318271	399558.0
6	0.08268023	4215783	0.3562352	2926215	537380.2	0.1658159	414744.7
7	0.09921628	4412471	0.3001825	3061849	650790.4	0.1465602	471881.0
8	0.11575233	4402258	0.2979999	3072796	642946.2	0.1424450	476584.6

```

9  0.13228838 4330241 0.3219284 3043569 758346.9  0.1567372 522090.4
10 0.14882442 4215395 0.3445074 2999347 701522.1  0.1542800 474190.3
11 0.16536047 4215395 0.3445074 2999347 701522.1  0.1542800 474190.3
12 0.18189652 4215395 0.3445074 2999347 701522.1  0.1542800 474190.3
13 0.19843256 4215395 0.3445074 2999347 701522.1  0.1542800 474190.3
14 0.21496861 4215395 0.3445074 2999347 701522.1  0.1542800 474190.3
15 0.23150466 4215395 0.3445074 2999347 701522.1  0.1542800 474190.3
16 0.24804070 4215395 0.3445074 2999347 701522.1  0.1542800 474190.3
17 0.26457675 4215395 0.3445074 2999347 701522.1  0.1542800 474190.3
18 0.28111280 4215395 0.3445074 2999347 701522.1  0.1542800 474190.3
19 0.29764884 4215395 0.3445074 2999347 701522.1  0.1542800 474190.3
20 0.31418489 4215395 0.3445074 2999347 701522.1  0.1542800 474190.3
21 0.33072094 4215395 0.3445074 2999347 701522.1  0.1542800 474190.3
22 0.34725699 4215395 0.3445074 2999347 701522.1  0.1542800 474190.3
23 0.36379303 4215395 0.3445074 2999347 701522.1  0.1542800 474190.3
24 0.38032908 4215395 0.3445074 2999347 701522.1  0.1542800 474190.3
25 0.39686513 4737126 0.2637852 3557586 738561.4  0.1501069 466476.1

```

```
cat("\nBest Tuned cp Value: ", caret_bp$bestTune$cp)
```

Best Tuned cp Value: 0.3803291

Variable Importance

```
caret::varImp(caret_bp)
```

rpart variable importance

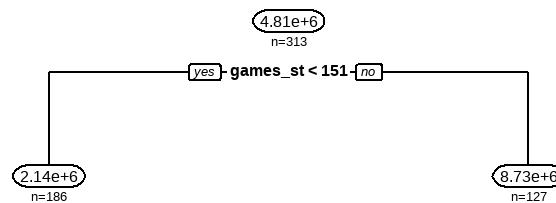
only 20 most important variables shown (out of 23)

	Overall
games_started	100.00
minutes_per_game	90.49
points	85.58
fg_made	83.88
ft_attempted	81.20
personal_fouls	0.00
fg_percent	0.00
turnovers	0.00
x3p_attempted	0.00
weight	0.00
blocks	0.00
height	0.00
x3p_made	0.00
fg_attempted	0.00
offensive_rebounds	0.00
steals	0.00
x3p_percent	0.00

ft_percent	0.00
games_played	0.00
ft_made	0.00

Show Best Tree (Caret)

```
prp(caret_bp$finalModel, type = 2, extra = 1, under = TRUE, digits = 3, fallen.leaves = TRUE)
```



```
cat("Number of Leaf Nodes in the Default Tree: ")
```

Number of Leaf Nodes in the Default Tree:

```
sum(caret_bp$finalModel$frame$var == "<leaf>")
```

[1] 2

Comments on Best Tree (Caret):

The CP value for Caret's best tree is lower than rparts best pruned tree, so caret's best tree will have more splits and thus more leaf nodes. Caret's best tree has 12 leaf nodes whereas rpart's best tree has 6 leaf nodes. The Variables that the models deem important are also different. Caret's top three most important variables are age, defensive rebounds, and fg attempted whereas rpart's top three most important variables

are games started, games played, and minutes per game. Due to the increase in splits there are more paths in caret's model all prompted by different predictor variables. For example both models top average salary is 20e+6 however the path of variables to get to that average salary is different.

For Rpart's model players who started 152 or more games, made at least 3.4 free throws, and attempted 6.7 or more free throws will have the average salary of about \$20 million.

For Caret's model players who started at least 185 games, made at least 2.85 free throws, are 33 years of age or older, weigh 233 pounds or more, and block at least 2.25 shots per game have an average salary of about \$20 million.

Evaluation metrics will be used to see if Caret's model performs better than rpart's model.

Evaluation (caret)

```
cbp_predict <- predict(caret_bp, testSet)
caret_bp_metrics <- round(forecast::accuracy(cbp_predict, testSet$salary),2)
rownames(caret_bp_metrics) <- "Caret Regression Tree"
print(caret_bp_metrics)
```

	ME	RMSE	MAE	MPE	MAPE
Caret Regression Tree	-474796.5	4585161	3254818	-295.88	321.03

```
rbind(caret_bp_metrics, bpt_metrics)
```

	ME	RMSE	MAE	MPE	MAPE
Caret Regression Tree	-474796.5	4585161	3254818	-295.88	321.03
Rpart Regression Tree	-150347.9	3210484	2345165	-231.71	256.97

Comments Comparing Evaluation across both regression trees:

The Rpart Regression Tree performs better than the Caret Regression Tree, due to its lower RMSE (3,372,027 vs. 4,283,480) and MAE (2,432,908 vs. 2,765,978), meaning its predictions are closer to the true salary values on average. Additionally, the Rpart model has smaller percentage errors suggesting it is more accurate and less biased.

Score (Caret)

```
caret_score <- predict(caret_bp, nba_score)
rf_pip <- cbind(nba_score, caret_score)
rf_pip
```

	player	age	height	weight	games_played	games_started	minutes_per_game	fg_made
1	Player1	36	81	260	984	820	31.8	6.8
2	Player2	32	76	185	714	504	30.2	5.4
3	Player3	27	85	235	636	472	32.0	5.2
	fg_attempted	fg_percent	x3p_made	x3p_attempted	x3p_percent	ft_made		
1	14.3	0.473	0.1	0.5	0.260	3.3		

```

2      11.0      0.437      1.6      4.2      0.384      5.0
3      11.8      0.438      1.0      3.0      0.339      2.6
  ft_attempted ft_percent offensive_rebounds defensive_rebounds assists blocks
1      4.4      0.765      3.1      6.3      1.8      0.3
2      5.7      0.870      0.6      2.6      2.1      0.1
3      3.4      0.786      1.0      3.9      1.7      0.6
  steals personal_fouls turnovers points caret_score
1      0.8      2.4      2.1    18.0    8725432
2      0.9      1.9      1.7    16.5    8725432
3      0.8      2.4      1.7    15.0    8725432

```

Data Partition (Random Forest)

```

p <- ncol(trainSet)-1

mtry_center <- max(1, round(p/3))
mtry_grid <- data.frame(mtry = sort(unique(pmax(1, c(mtry_center-1, mtry_center, mtry_center+1,
2, p)))))

```

Modeling (Random Forest)

```

set.seed(1)
rf_model <- train(salary ~ .,
                    data = trainSet,
                    method = "rf",
                    trControl = myCtrl,
                    tuneGrid = mtry_grid,
                    ntree = 1000,
                    importance = TRUE)

```

Resample

```
rf_model
```

Random Forest

313 samples
23 predictor

No pre-processing
Resampling: Cross-Validated (10 fold)
Summary of sample sizes: 282, 281, 281, 282, 282, 282, ...
Resampling results across tuning parameters:

mtry	RMSE	Rsquared	MAE

```

2    3472405  0.5658555  2389507
7    3432674  0.5766471  2333652
8    3408193  0.5865120  2310666
9    3393808  0.5871320  2297864
23   3441298  0.5746062  2317649

```

RMSE was used to select the optimal model using the smallest value.
The final value used for the model was mtry = 9.

Variable Importance

```
caret::varImp(rf_model)
```

rf variable importance

only 20 most important variables shown (out of 23)

	Overall
games_started	100.000
games_played	94.822
minutes_per_game	65.573
age	60.315
points	56.444
fg_made	44.587
defensive_rebounds	39.359
ft_made	31.713
offensive_rebounds	31.387
fg_percent	31.167
ft_attempted	30.327
fg_attempted	28.804
personal_fouls	17.223
steals	15.845
turnovers	15.300
assists	13.600
x3p_attempted	12.898
x3p_percent	8.699
blocks	6.380
ft_percent	2.828

Comments on Random Forest Variable Importance:

The top four variables by variable importance is games started, games played, age, and minutes per game.

Evaluation (Random Forest)

```

predicted_rf <- predict(rf_model, testSet)
test_rf <- round(forecast::accuracy(predicted_rf, testSet$salary),2)

```

```
rownames(test_rf) <- "Random Forest"
test_rf
```

	ME	RMSE	MAE	MPE	MAPE
Random Forest	-292188.9	3259179	2243796	-224.08	246.02

Evaluation Metrics Across All Trees

```
rbind(test_rf, caret_bp_metrics, bpt_metrics)
```

	ME	RMSE	MAE	MPE	MAPE
Random Forest	-292188.9	3259179	2243796	-224.08	246.02
Caret Regression Tree	-474796.5	4585161	3254818	-295.88	321.03
Rpart Regression Tree	-150347.9	3210484	2345165	-231.71	256.97

Comments on Evaluation Metrics:

The Rpart Regression Tree performs best overall, with the lowest RMSE (3,372,027) and MAE (2,432,908), indicating the smallest average and squared prediction errors among the models. It also has the lowest MAPE (158.93%), meaning its salary predictions are the closest to the true values. While the Random Forest performs similarly, its higher percentage errors suggest slightly less consistent prediction accuracy.

Score (Random Forest)

```
score_rf <- predict(rf_model, nba_score)
rf_pip <- cbind(nba_score, score_rf, caret_score, rpart_score)
rf_pip
```

	player	age	height	weight	games_played	games_started	minutes_per_game	fg_made
1	Player1	36	81	260	984	820	31.8	6.8
2	Player2	32	76	185	714	504	30.2	5.4
3	Player3	27	85	235	636	472	32.0	5.2
	fg_attempted	fg_percent	x3p_made	x3p_attempted	x3p_percent	ft_made		
1	14.3	0.473	0.1	0.5	0.260	3.3		
2	11.0	0.437	1.6	4.2	0.384	5.0		
3	11.8	0.438	1.0	3.0	0.339	2.6		
	ft_attempted	ft_percent	offensive_rebounds	defensive_rebounds	assists	blocks		
1	4.4	0.765	3.1	6.3	1.8	0.3		
2	5.7	0.870	0.6	2.6	2.1	0.1		
3	3.4	0.786	1.0	3.9	1.7	0.6		
	steals	personal_fouls	turnovers	points	score_rf	caret_score	rpart_score	
1	0.8	2.4	2.1	18.0	7240941	8725432	3454649	
2	0.9	1.9	1.7	16.5	11604523	8725432	13977446	
3	0.8	2.4	1.7	15.0	10345555	8725432	7843275	

Comments on Scoring Across Tree Models:

We trust the Rpart score predictions more because the evaluation metrics clearly show that the Rpart

Regression Tree achieved the lowest RMSE (3,372,027) and lowest MAE (2,432,908) among all models. These values indicate that its predictions are, on average, closest to the true player salaries, with smaller overall errors and less variation.